

Deep Learning

Lecture 3: Convolutional networks

Prof. Gilles Louppe
g.louppe@uliege.be



Today

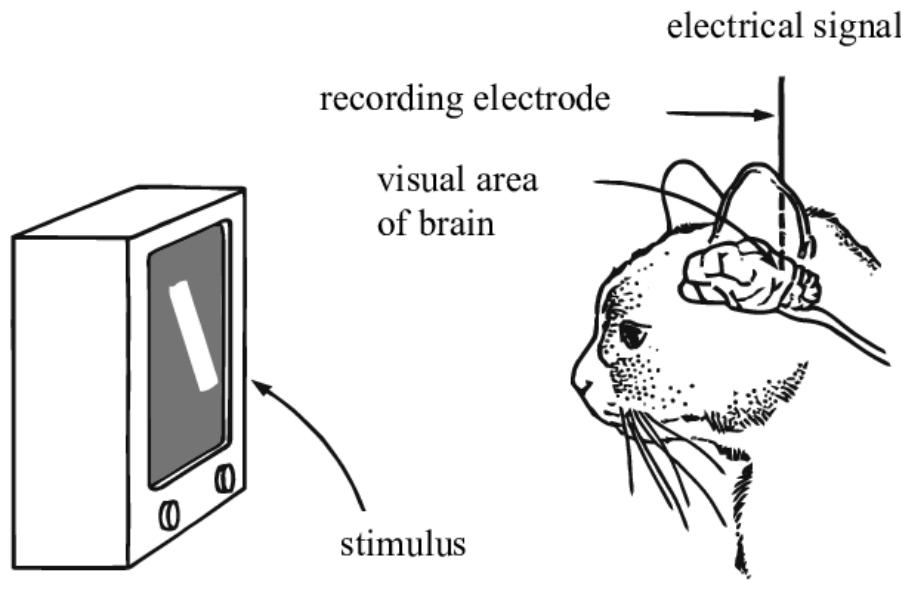
How to make neural networks see?

- A little history
- Convolutions
- Pooling
- Convolutional networks
- Under the hood

A little history

Visual perception (Hubel and Wiesel, 1959-1962)

- David Hubel and Torsten Wiesel discover the neural basis of **visual perception**.
- Awarded the Nobel Prize of Medicine in 1981 for their discovery.





Hubel and Wiesel Cat Experiment



Watch later

Share



Hubel and Wiesel



Hubel & Wiesel 1: Intro



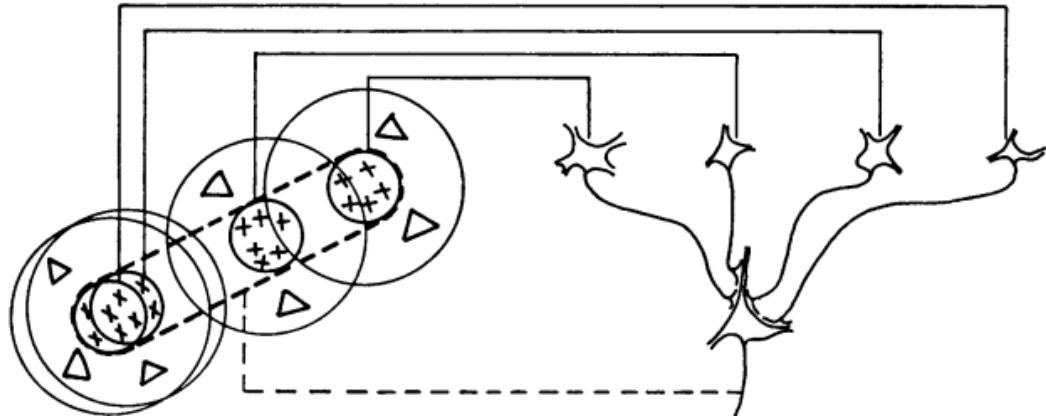
Watch later



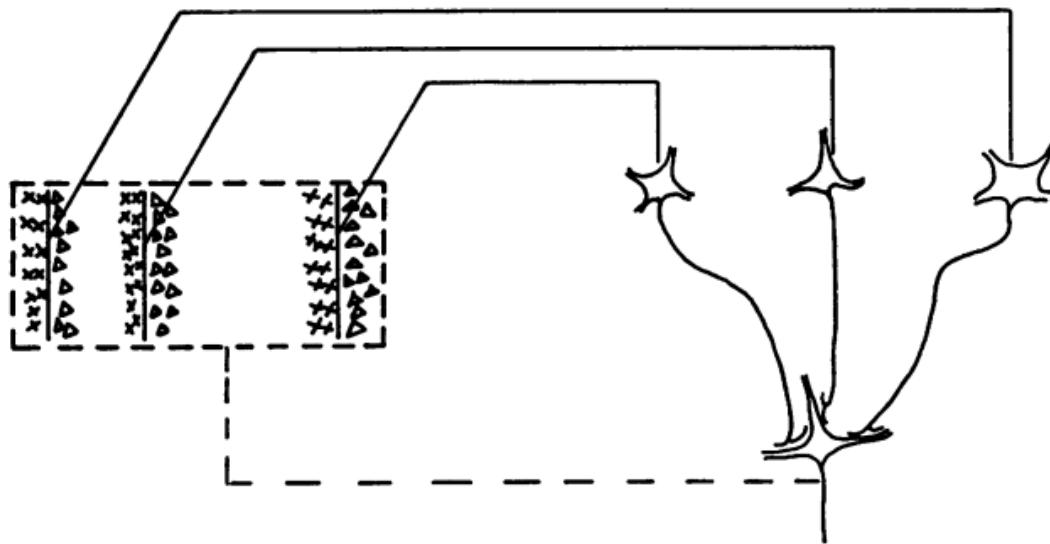
Share



Hubel and Wiesel

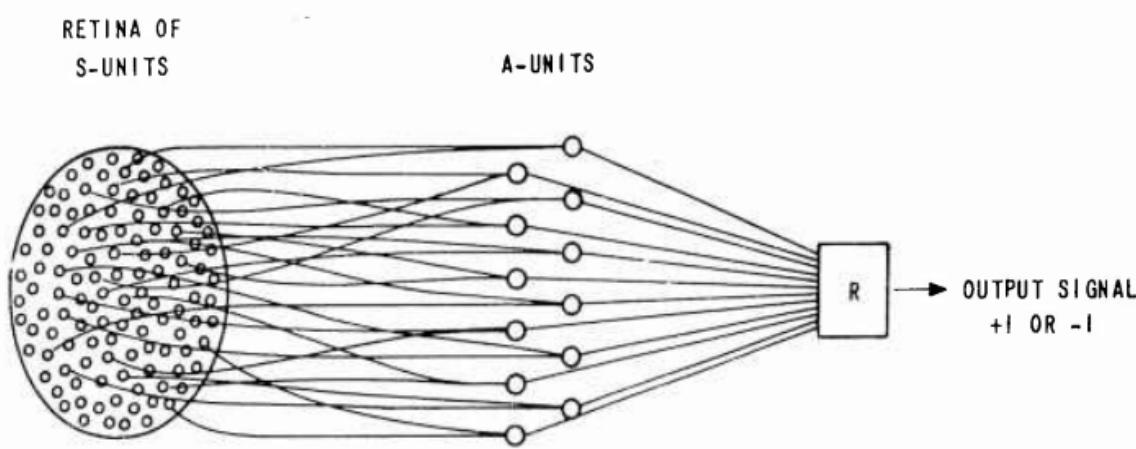


Text-fig. 19. Possible scheme for explaining the organization of simple receptive fields. A large number of lateral geniculate cells, of which four are illustrated in the upper right in the figure, have receptive fields with 'on' centres arranged along a straight line on the retina. All of these project upon a single cortical cell, and the synapses are supposed to be excitatory. The receptive field of the cortical cell will then have an elongated 'on' centre indicated by the interrupted lines in the receptive-field diagram to the left of the figure.

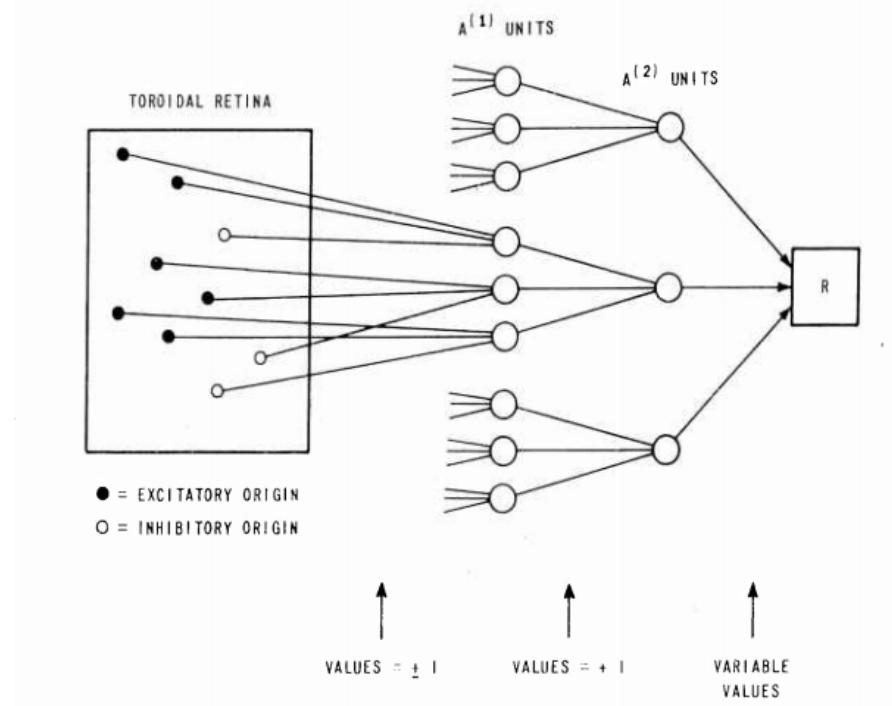


Text-fig. 20. Possible scheme for explaining the organization of complex receptive fields. A number of cells with simple fields, of which three are shown schematically, are imagined to project to a single cortical cell of higher order. Each projecting neurone has a receptive field arranged as shown to the left: an excitatory region to the left and an inhibitory region to the right of a vertical straight-line boundary. The boundaries of the fields are staggered within an area outlined by the interrupted lines. Any vertical-edge stimulus falling across this rectangle, regardless of its position, will excite some simple-field cells, leading to excitation of the higher-order cell.

The Mark-1 Perceptron (Rosenblatt, 1957-61)



- Rosenblatt builds the first implementation of a neural network.
- The network is an analog circuit. Parameters are potentiometers.



"If we show the perceptron a stimulus, say a square, and associate a response to that square, this response will immediately generalize perfectly to all transforms of the square under the transformation group [...]."

AI winter (Minsky and Papert, 1969+)

- Minsky and Papert prove a series of impossibility results for the perceptron (or rather, a narrowly defined variant thereof).
- AI winter follows.

Theorem 0.8: No diameter-limited perceptron can determine whether or not all the parts of any geometric figure are connected to one another! That is, no such perceptron computes $\psi_{\text{CONNECTED}}$.



Actual Variable	Variable Number (Address)	Category	Major Source	Minor Source
$A(t+1)$	13	sum	12	11
$k_1 A(t)$	12	product	3	1
$k_2^* U(t) \left(\frac{A(t)-U(t)}{A(t)+U(t)}\right)^{k_4}$	11	product	10	4
$U(t) \left(\frac{A(t)-U(t)}{A(t)+U(t)}\right)^{k_4}$	10	product	9	2
$\left(\frac{A(t)-U(t)}{A(t)+U(t)}\right)^{k_4}$	9	power	8	5
$\frac{A(t)-U(t)}{A(t)+U(t)}$	8	ratio	7	6
$A(t)-U(t)$	7	difference	1	2
$A(t)+U(t)$	6	sum	1	2
k_4	5	parameter	-	-
k_2	4	parameter	-	-
k_1	3	parameter	-	-
$U(t)$	2	given	-	-
$A(t)$	1	given	-	-

Automatic differentiation (Werbos, 1974)

- Werbos formulate an arbitrary function as a computational graph.
- Symbolic derivatives are computed by dynamic programming.

Neocognitron (Fukushima, 1980)

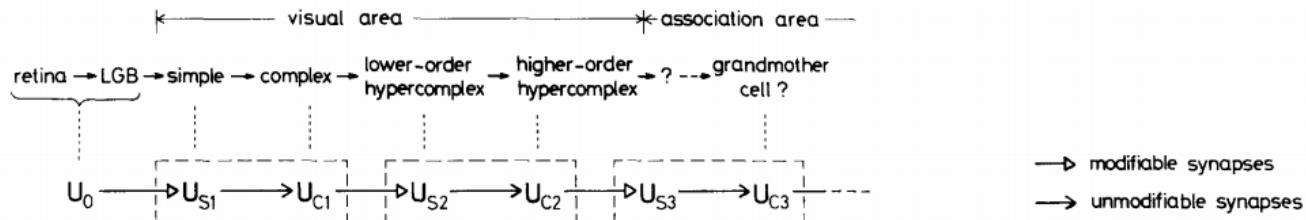


Fig. 1. Correspondence between the hierarchy model by Hubel and Wiesel, and the neural network of the neocognitron

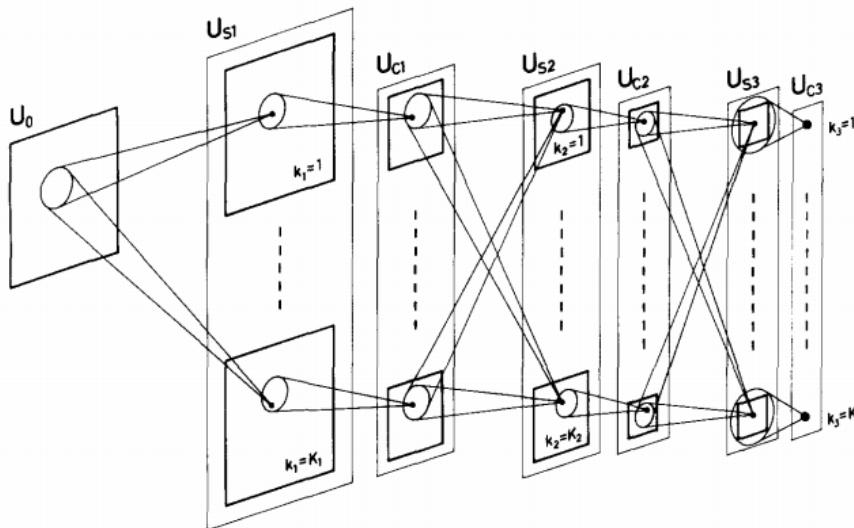
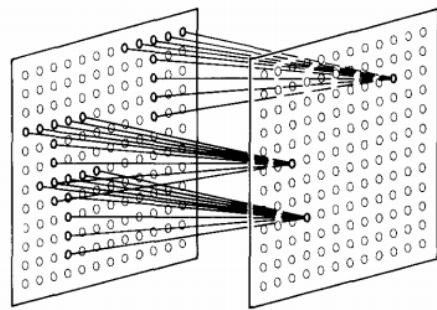
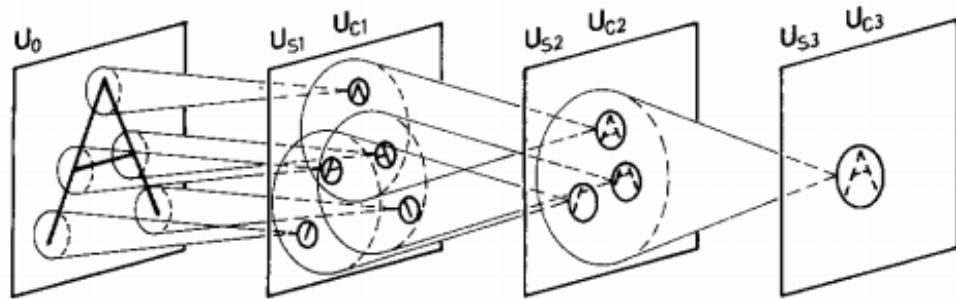


Fig. 2. Schematic diagram illustrating the interconnections between layers in the neocognitron

Fukushima proposes a direct neural network implementation of the hierarchy model of the visual nervous system of Hubel and Wiesel.



Convolutions



Feature hierarchy

Backpropagation (Rumelhart et al, 1986)

- Rumelhart and Hinton introduce **backpropagation** in multi-layer networks with sigmoid non-linearities and sum of squares loss function.
- They advocate for batch gradient descent in supervised learning.
- Discuss online gradient descent, momentum and random initialization.
- Depart from **biologically plausible** training algorithms.

The backward pass starts by computing $\partial E / \partial y$ for each of the output units. Differentiating equation (3) for a particular case, c , and suppressing the index c gives

$$\partial E / \partial y_j = y_j - d_j \quad (4)$$

We can then apply the chain rule to compute $\partial E / \partial x_j$,

$$\partial E / \partial x_j = \partial E / \partial y_j \cdot dy_j / dx_j$$

Differentiating equation (2) to get the value of dy_j / dx_j and substituting gives

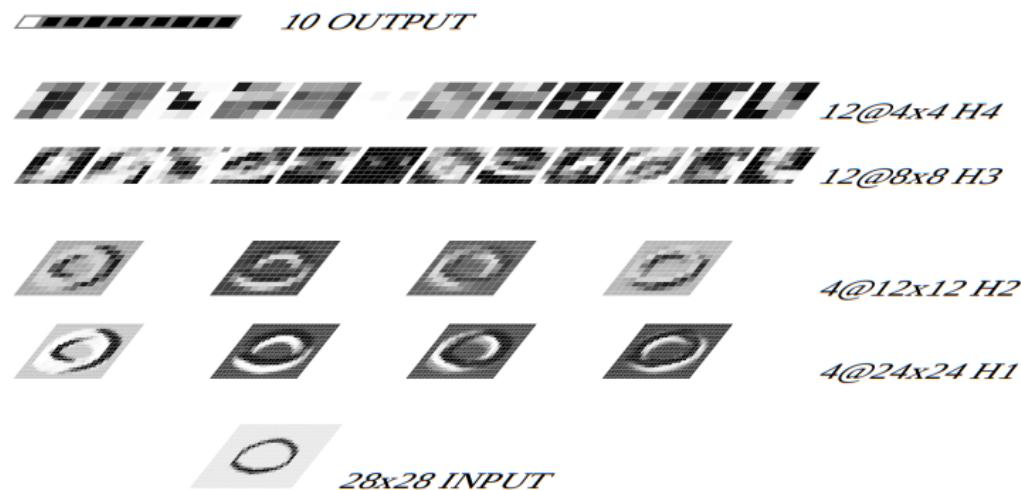
$$\partial E / \partial x_j = \partial E / \partial y_j \cdot y_j(1 - y_j) \quad (5)$$

This means that we know how a change in the total input x to an output unit will affect the error. But this total input is just a linear function of the states of the lower level units and it is also a linear function of the weights on the connections, so it is easy to compute how the error will be affected by changing these states and weights. For a weight w_{ji} , from i to j the derivative is

$$\begin{aligned} \partial E / \partial w_{ji} &= \partial E / \partial x_j \cdot \partial x_j / \partial w_{ji} \\ &= \partial E / \partial x_j \cdot y_i \end{aligned} \quad (6)$$

Convolutional networks (LeCun, 1990)

- LeCun trains a convolutional network by backpropagation.
- He advocates for end-to-end feature learning in image classification.





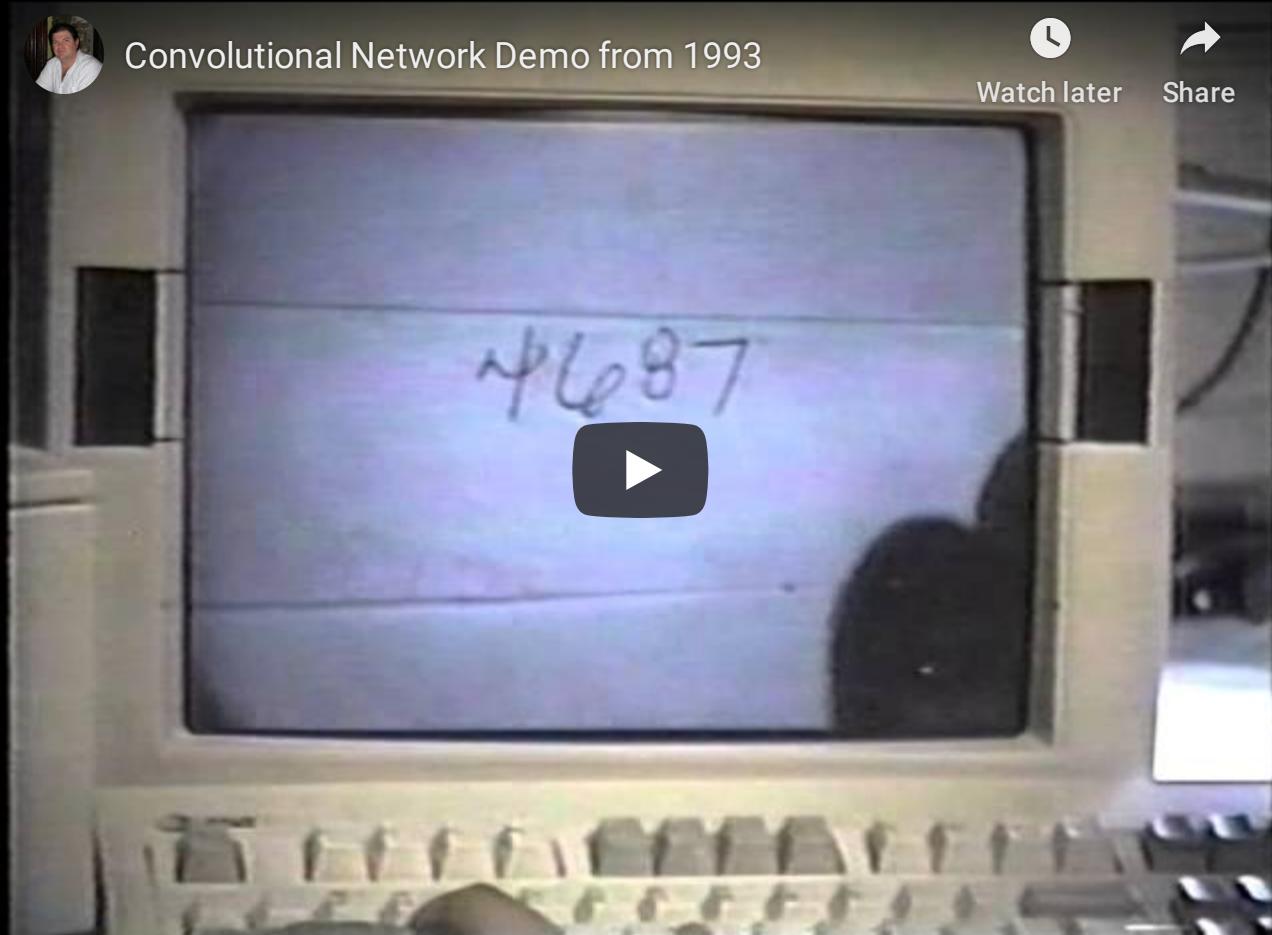
Convolutional Network Demo from 1993



Watch later



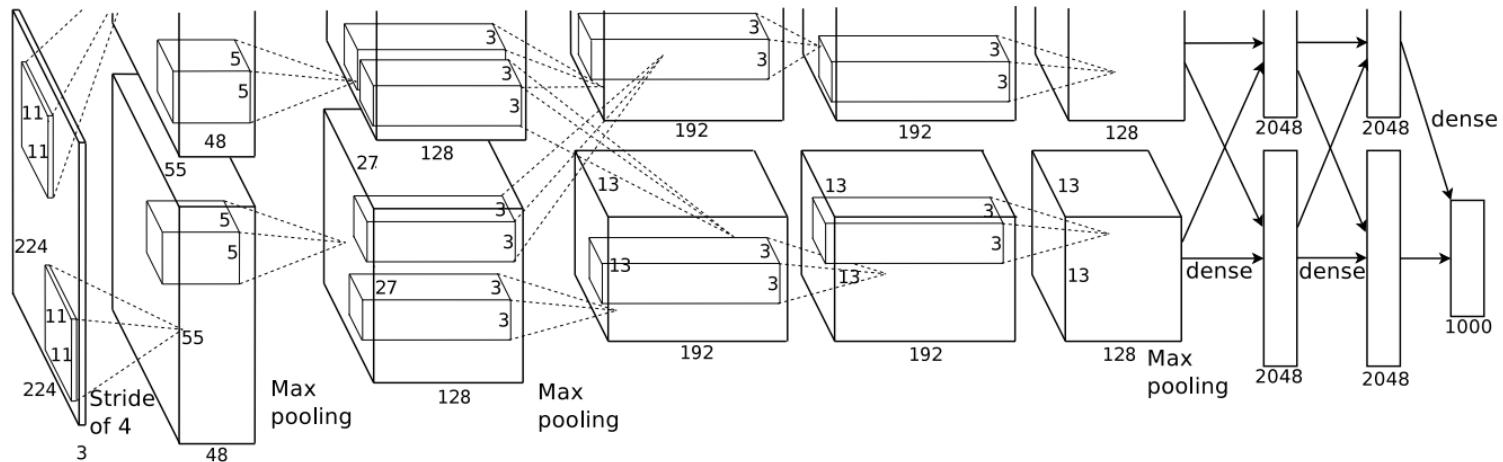
Share



LeNet-1 (LeCun et al, 1993)

AlexNet (Krizhevsky et al, 2012)

- Krizhevsky trains a convolutional network on ImageNet with two GPUs.
- 16.4% top-5 error on ILSVRC'12, outperforming all other entries by 10% or more.
- This event triggers the deep learning revolution.



Convolutions

If they were handled as normal "unstructured" vectors, large-dimension signal as sound samples or images would require of intractable size.

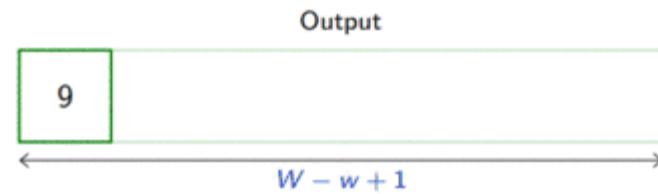
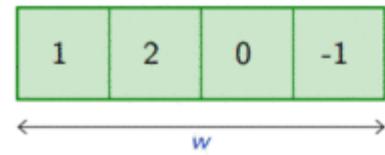
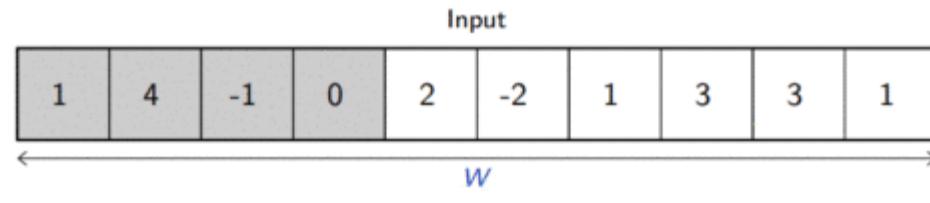
E.g., a linear layer taking 256×256 RGB images as input and producing an image of same size would require

$$(256 \times 256 \times 3)^2 \approx 3.87e + 10$$

parameters, with the corresponding memory footprint (150Gb!), and excess of capacity.

This requirement is also inconsistent with the intuition that such large signals have some "invariance in translation". **A representation meaningful at a certain location can / should be used everywhere.**

A convolution layer embodies this idea. It applies the same linear transformation locally everywhere while preserving the signal structure.



Convolutions

For one-dimensional tensors, given an input vector $\mathbf{x} \in \mathbb{R}^W$ and a convolutional kernel $\mathbf{u} \in \mathbb{R}^w$, the discrete **convolution** $\mathbf{x} \circledast \mathbf{u}$ is a vector of size $W - w + 1$ such that

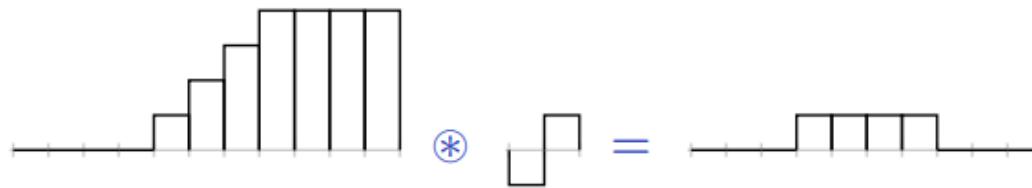
$$(\mathbf{x} \circledast \mathbf{u})[i] = \sum_{m=0}^{w-1} x_{m+i} u_m.$$

Note

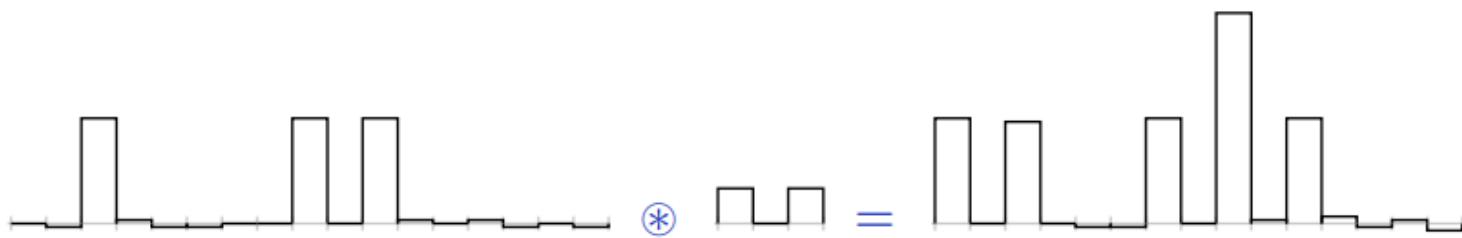
Technically, \circledast denotes the cross-correlation operator. However, most machine learning libraries call it convolution.

Convolutions can implement differential operators:

$$(0, 0, 0, 0, 1, 2, 3, 4, 4, 4, 4) \circledast (-1, 1) = (0, 0, 0, 1, 1, 1, 1, 0, 0, 0)$$

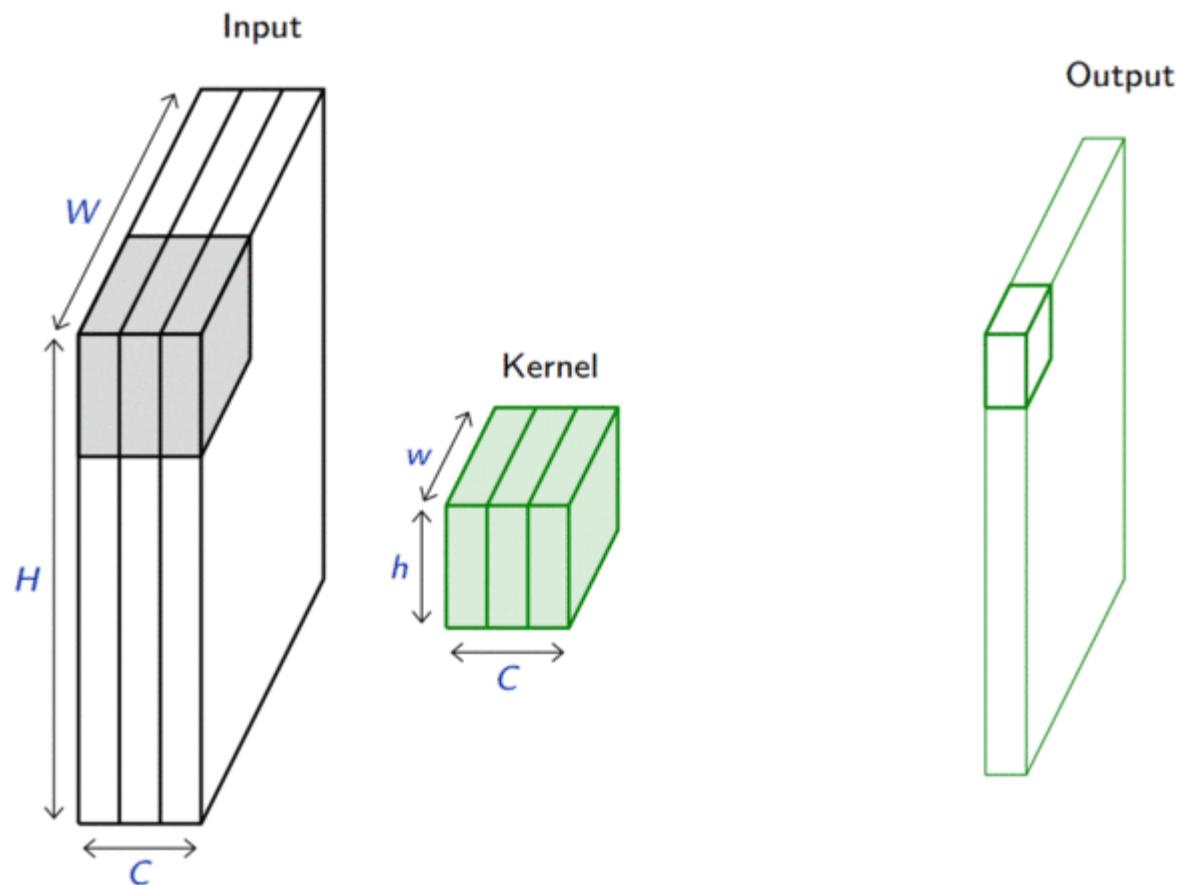


or crude template matchers:



Convolutions generalize to multi-dimensional tensors:

- In its most usual form, a convolution takes as input a 3D tensor $\mathbf{x} \in \mathbb{R}^{C \times H \times W}$, called the **input feature map**.
- A kernel $\mathbf{u} \in \mathbb{R}^{C \times h \times w}$ slides across the input feature map, along its height and width. The size $h \times w$ is the size of the **receptive field**.
- At each location, the element-wise product between the kernel and the input elements it overlaps is computed and the results are summed up.



- The final output \mathbf{o} is a 2D tensor of size $(H - h + 1) \times (W - w + 1)$ called the **output feature map** and such that:

$$\mathbf{o}_{j,i} = \mathbf{b}_{j,i} + \sum_{c=0}^{C-1} (\mathbf{x}_c \circledast \mathbf{u}_c)[j, i] = \mathbf{b}_{j,i} + \sum_{c=0}^{C-1} \sum_{n=0}^{h-1} \sum_{m=0}^{w-1} \mathbf{x}_{c,n+j,m+i} \mathbf{u}_{c,n,m}$$

where \mathbf{u} and \mathbf{b} are shared parameters to learn.

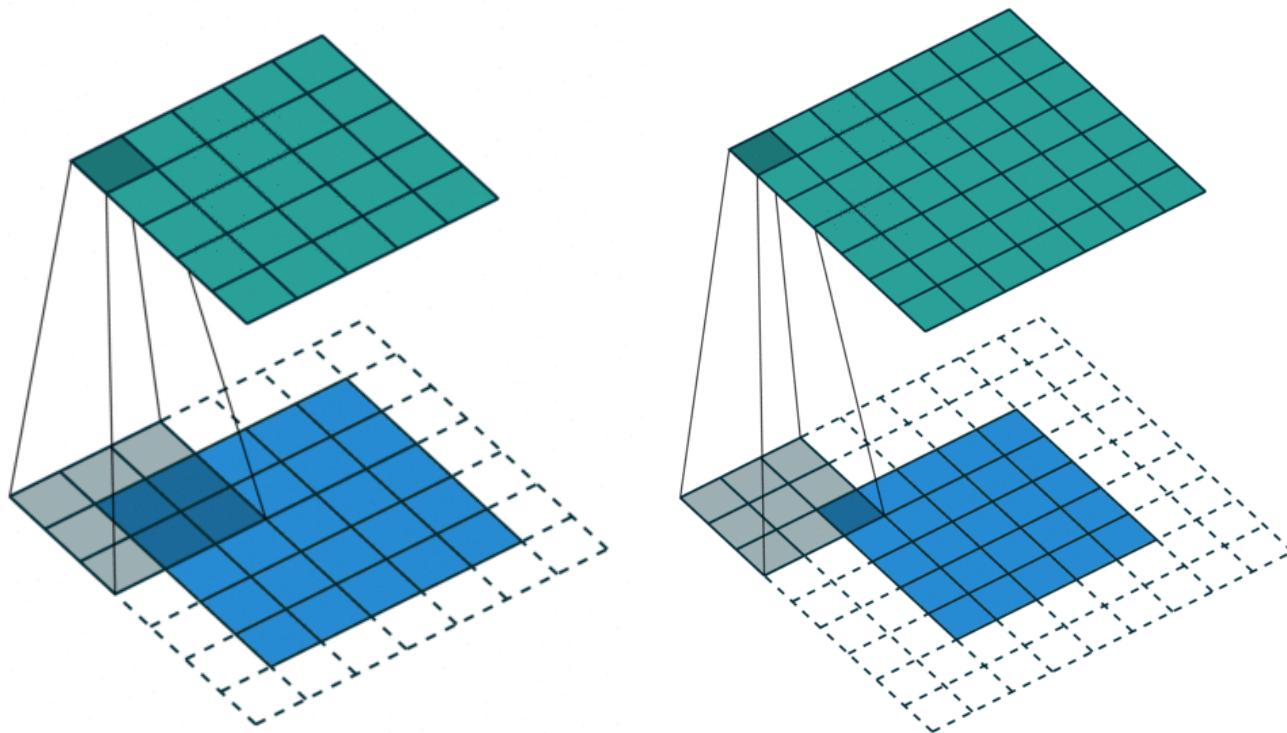
- D convolutions can be applied in the same way to produce a $D \times (H - h + 1) \times (W - w + 1)$ feature map, where D is the depth.
- Swiping across channels with a 3D convolution usually makes no sense, unless the channel index has some metric meaning.

Convolutions have three additional parameters:

- The **padding** specifies the size of a zeroed frame added around the input.
- The **stride** specifies a step size when moving the kernel across the signal.
- The **dilation** modulates the expansion of the filter without adding weights.

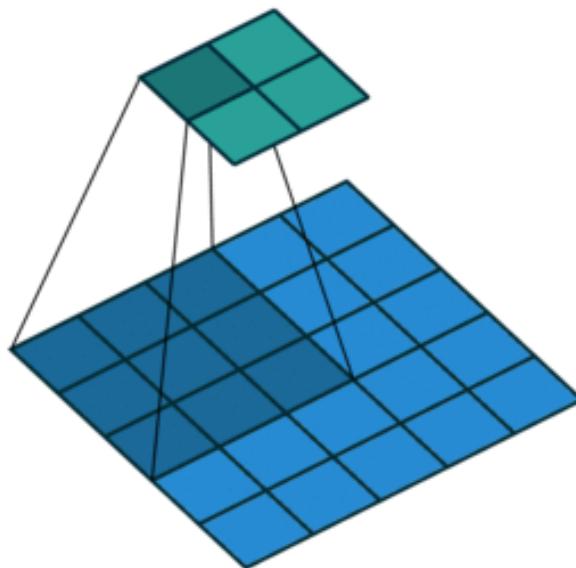
Padding

Padding is useful to control the spatial dimension of the feature map, for example to keep it constant across layers.



Strides

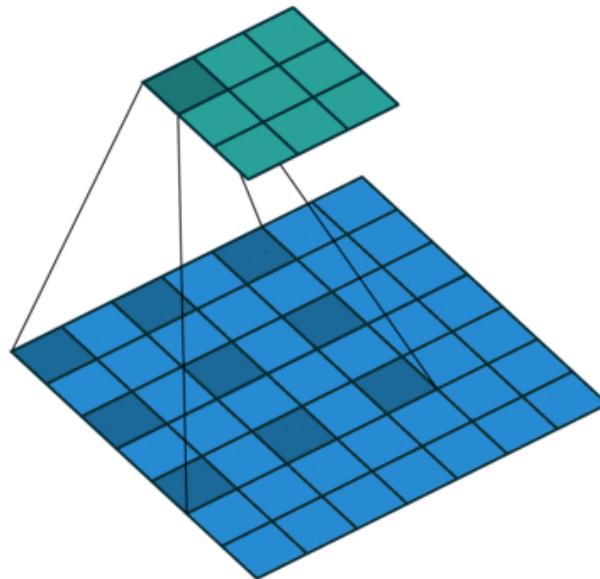
Stride is useful to reduce the spatial dimension of the feature map by a constant factor.



Dilation

The dilation modulates the expansion of the kernel support by adding rows and columns of zeros between coefficients.

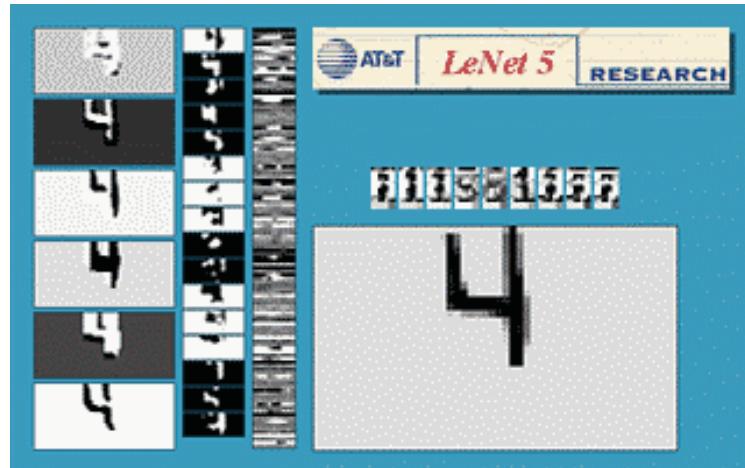
Having a dilation coefficient greater than one increases the units receptive field size without increasing the number of parameters.



Equivariance

A function f is **equivariant** to g if $f(g(\mathbf{x})) = g(f(\mathbf{x}))$.

- Parameter sharing used in a convolutional layer causes the layer to be equivariant to translation.
- That is, if g is any function that translates the input, the convolution function is equivariant to g .



If an object moves in the input image, its representation will move the same amount in the output.

- Equivariance is useful when we know some local function is useful everywhere (e.g., edge detectors).
- Convolution is not equivariant to other operations such as change in scale or rotation.

Convolutions as matrix multiplications

As a guiding example, let us consider the convolution of single-channel tensors $\mathbf{x} \in \mathbb{R}^{4 \times 4}$ and $\mathbf{u} \in \mathbb{R}^{3 \times 3}$:

$$\mathbf{x} \circledast \mathbf{u} = \begin{pmatrix} 4 & 5 & 8 & 7 \\ 1 & 8 & 8 & 8 \\ 3 & 6 & 6 & 4 \\ 6 & 5 & 7 & 8 \end{pmatrix} \circledast \begin{pmatrix} 1 & 4 & 1 \\ 1 & 4 & 3 \\ 3 & 3 & 1 \end{pmatrix} = \begin{pmatrix} 122 & 148 \\ 126 & 134 \end{pmatrix}$$

The convolution operation can be equivalently re-expressed as a single matrix multiplication:

- the convolutional kernel \mathbf{u} is rearranged as a **sparse Toeplitz circulant matrix**, called the convolution matrix:

$$\mathbf{U} = \begin{pmatrix} 1 & 4 & 1 & 0 & 1 & 4 & 3 & 0 & 3 & 3 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 4 & 1 & 0 & 1 & 4 & 3 & 0 & 3 & 3 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 4 & 1 & 0 & 1 & 4 & 3 & 0 & 3 & 3 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 4 & 1 & 0 & 1 & 4 & 3 & 0 & 3 & 3 & 1 \end{pmatrix}$$

- the input \mathbf{x} is flattened row by row, from top to bottom:

$$v(\mathbf{x}) = (4 \quad 5 \quad 8 \quad 7 \quad 1 \quad 8 \quad 8 \quad 8 \quad 3 \quad 6 \quad 6 \quad 4 \quad 6 \quad 5 \quad 7 \quad 8)^T$$

Then,

$$\mathbf{U}v(\mathbf{x}) = (122 \quad 148 \quad 126 \quad 134)^T$$

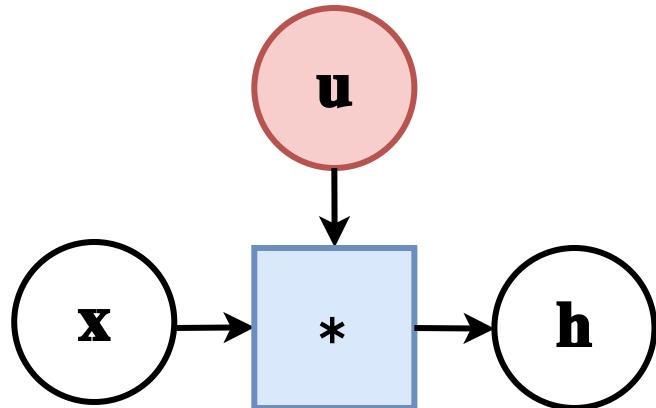
which we can reshape to a 2×2 matrix to obtain $\mathbf{x} \circledast \mathbf{u}$.

The same procedure generalizes to $\mathbf{x} \in \mathbb{R}^{H \times W}$ and convolutional kernel $\mathbf{u} \in \mathbb{R}^{h \times w}$, such that:

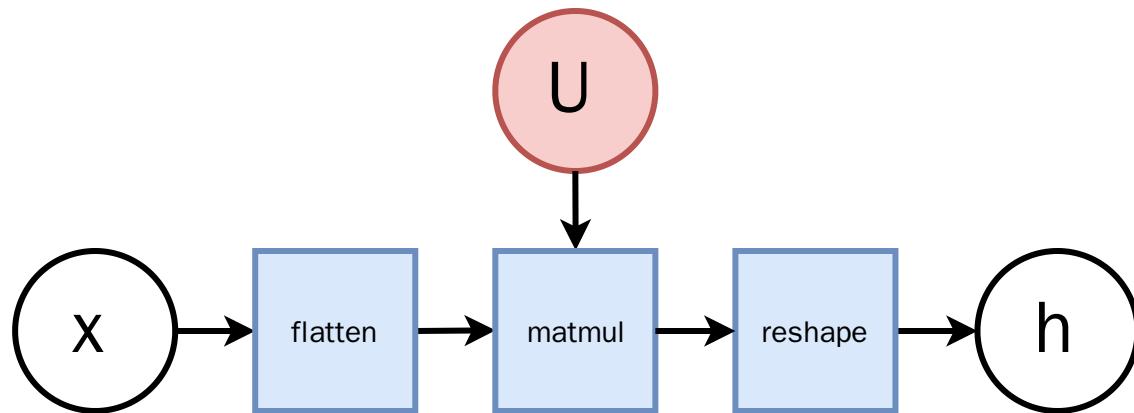
- the convolutional kernel is rearranged as a sparse Toeplitz circulant matrix \mathbf{U} of shape $(H - h + 1)(W - w + 1) \times HW$ where
 - each row i identifies an element of the output feature map,
 - each column j identifies an element of the input feature map,
 - the value $\mathbf{U}_{i,j}$ corresponds to the kernel value the element j is multiplied with in output i ;
- the input \mathbf{x} is flattened into a column vector $v(\mathbf{x})$ of shape $HW \times 1$;
- the output feature map $\mathbf{x} \circledast \mathbf{u}$ is obtained by reshaping the $(H - h + 1)(W - w + 1) \times 1$ column vector $\mathbf{U}v(\mathbf{x})$ as a $(H - h + 1) \times (W - w + 1)$ matrix.

Therefore, a convolutional layer is a special case of a fully connected layer:

$$\mathbf{h} = \mathbf{x} \circledast \mathbf{u} \Leftrightarrow v(\mathbf{h}) = \mathbf{U}v(\mathbf{x}) \Leftrightarrow v(\mathbf{h}) = \mathbf{W}^T v(\mathbf{x})$$



\Leftrightarrow



Pooling

When the input volume is large, **pooling layers** can be used to reduce the input dimension while preserving its global structure, in a way similar to a down-scaling operation.

Pooling

Consider a pooling area of size $h \times w$ and a 3D input tensor $\mathbf{x} \in \mathbb{R}^{C \times (rh) \times (sw)}$.

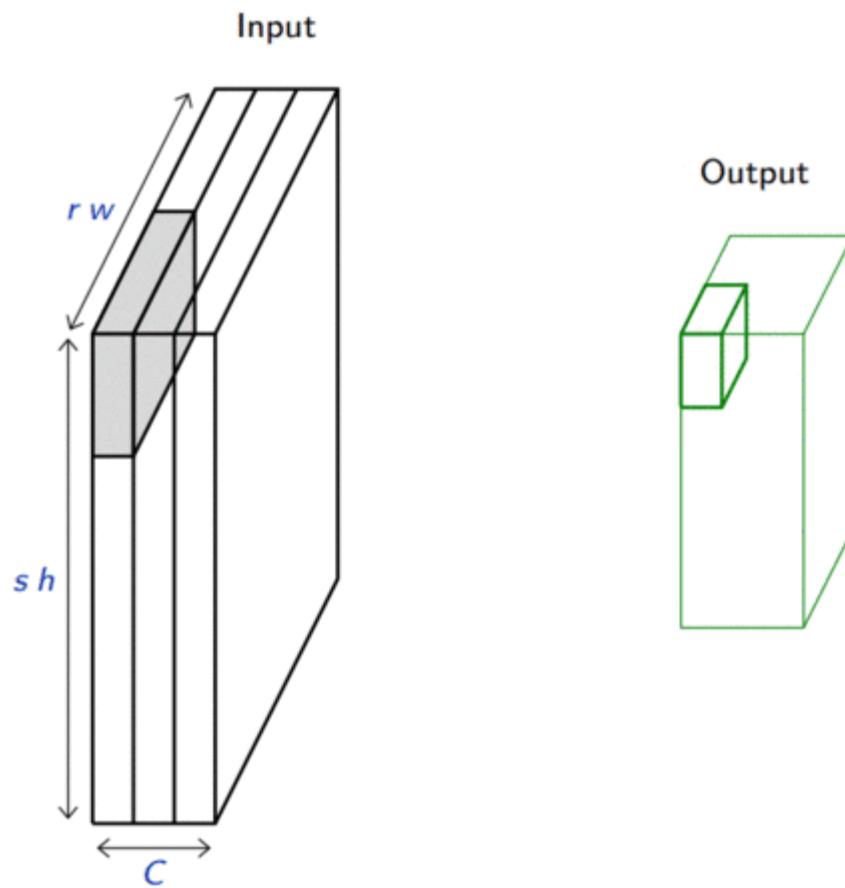
- Max-pooling produces a tensor $\mathbf{o} \in \mathbb{R}^{C \times r \times s}$ such that

$$\mathbf{o}_{c,j,i} = \max_{n < h, m < w} \mathbf{x}_{c,rj+n,si+m}.$$

- Average pooling produces a tensor $\mathbf{o} \in \mathbb{R}^{C \times r \times s}$ such that

$$\mathbf{o}_{c,j,i} = \frac{1}{hw} \sum_{n=0}^{h-1} \sum_{m=0}^{w-1} \mathbf{x}_{c,rj+n,si+m}.$$

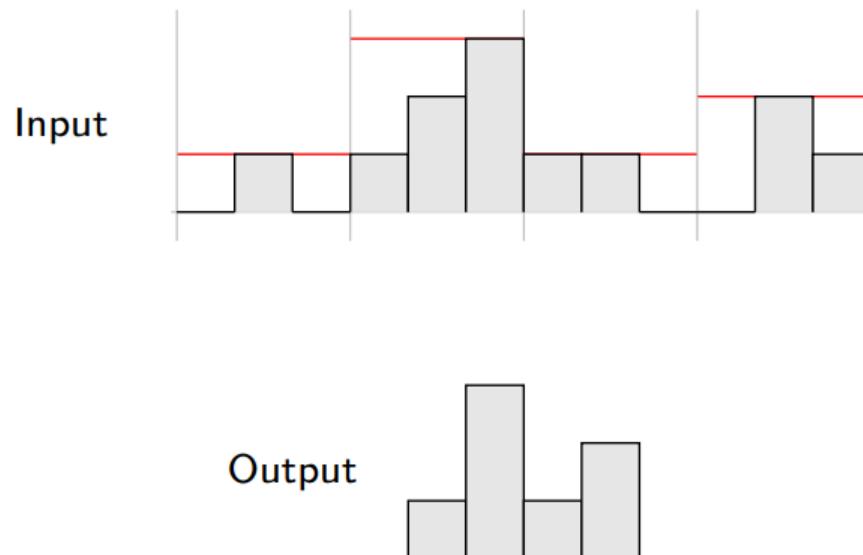
Pooling is very similar in its formulation to convolution.



Invariance

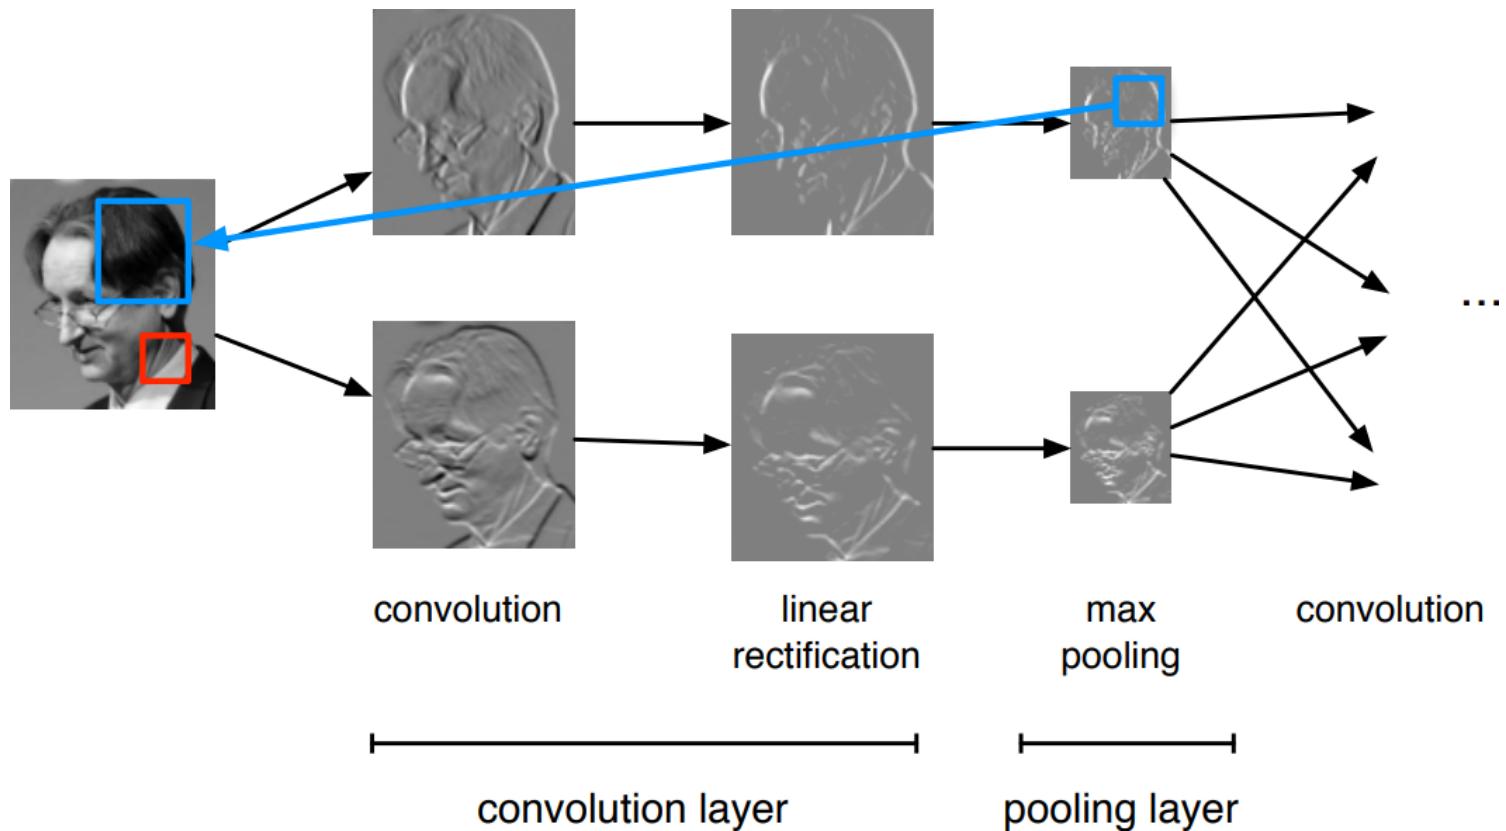
A function f is **invariant** to g if $f(g(\mathbf{x})) = f(\mathbf{x})$.

- Pooling layers provide invariance to any permutation inside one cell.
- It results in (pseudo-)invariance to local translations.
- This helpful if we care more about the presence of a pattern rather than its exact position.



Convolutional networks

A **convolutional network** is generically defined as a composition of convolutional layers (**CONV**), pooling layers (**POOL**), linear rectifiers (**RELU**) and fully connected layers (**FC**).



The most common convolutional network architecture follows the pattern:

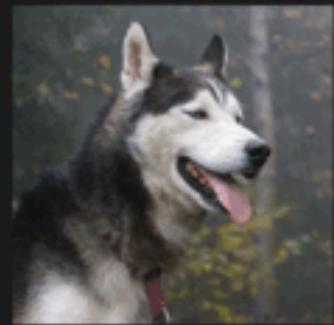
INPUT → [[**CONV** → **RELU**]* N → **POOL?**]* M → [**FC** → **RELU**]* K → **FC**

where:

- * indicates repetition;
- **POOL?** indicates an optional pooling layer;
- $N \geq 0$ (and usually $N \leq 3$), $M \geq 0$, $K \geq 0$ (and usually $K < 3$);
- the last fully connected layer holds the output (e.g., the class scores).

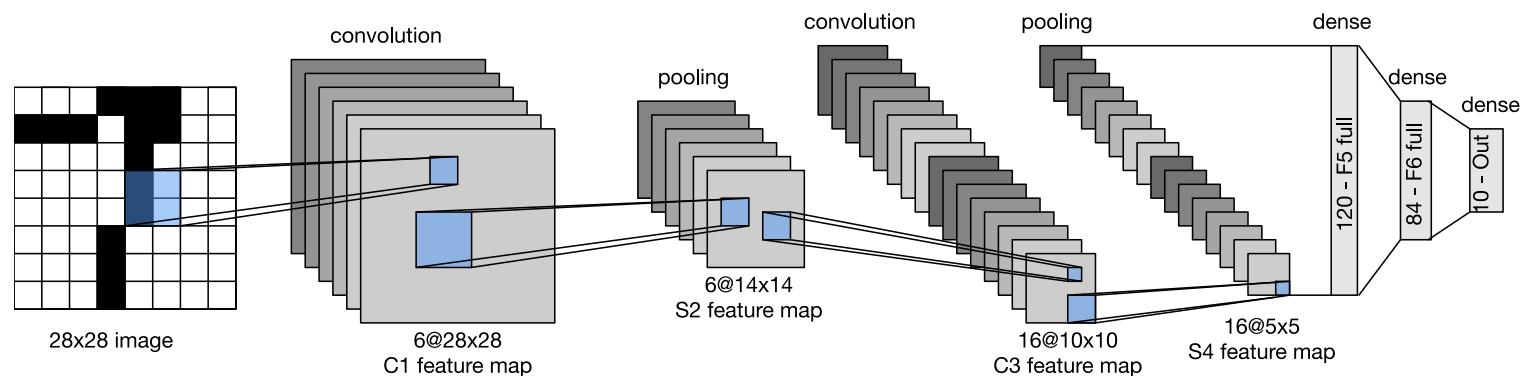
Some common architectures for convolutional networks following this pattern include:

- INPUT → FC, which implements a linear classifier ($N = M = K = 0$).
- INPUT → [FC → RELU] $*K$ → FC, which implements a K -layer MLP.
- INPUT → CONV → RELU → FC.
- INPUT → [CONV → RELU → POOL] $*2$ → FC → RELU → FC.
- INPUT → [[CONV → RELU] $*2$ → POOL] $*3$ → [FC → RELU] $*2$ → FC.



LeNet-5 (LeCun et al, 1998)

Composition of two **CONV + POOL** layers, followed by a block of fully-connected layers.

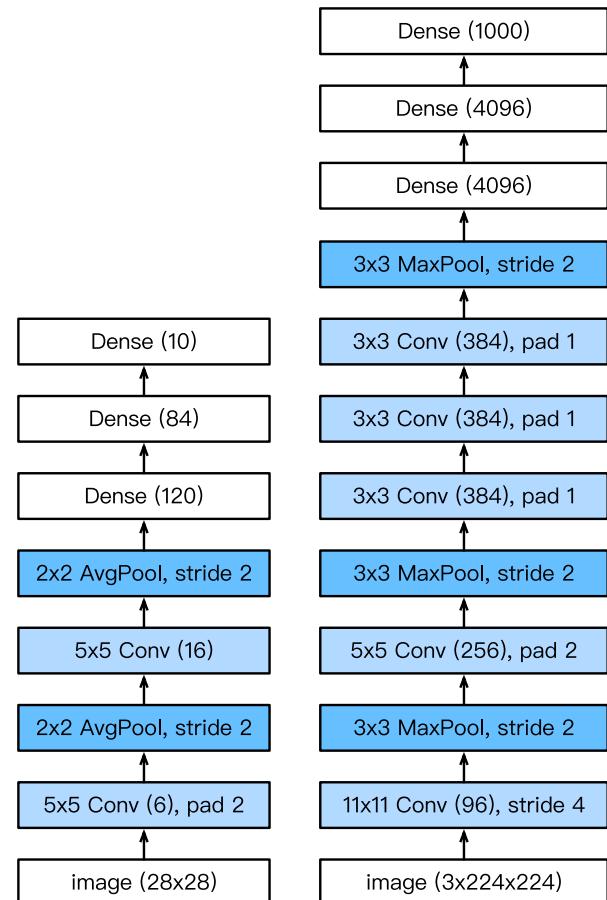


Layer (type)	Output Shape	Param #
<hr/>		
Conv2d-1	[-1, 6, 28, 28]	156
ReLU-2	[-1, 6, 28, 28]	0
MaxPool2d-3	[-1, 6, 14, 14]	0
Conv2d-4	[-1, 16, 10, 10]	2,416
ReLU-5	[-1, 16, 10, 10]	0
MaxPool2d-6	[-1, 16, 5, 5]	0
Conv2d-7	[-1, 120, 1, 1]	48,120
ReLU-8	[-1, 120, 1, 1]	0
Linear-9	[-1, 84]	10,164
ReLU-10	[-1, 84]	0
Linear-11	[-1, 10]	850
LogSoftmax-12	[-1, 10]	0
<hr/>		
Total params:	61,706	
Trainable params:	61,706	
Non-trainable params:	0	
<hr/>		
Input size (MB):	0.00	
Forward/backward pass size (MB):	0.11	
Params size (MB):	0.24	
Estimated Total Size (MB):	0.35	
<hr/>		

AlexNet (Krizhevsky et al, 2012)

Composition of a 8-layer convolutional neural network with a 3-layer MLP.

The original implementation was made of two parts such that it could fit within two GPUs.

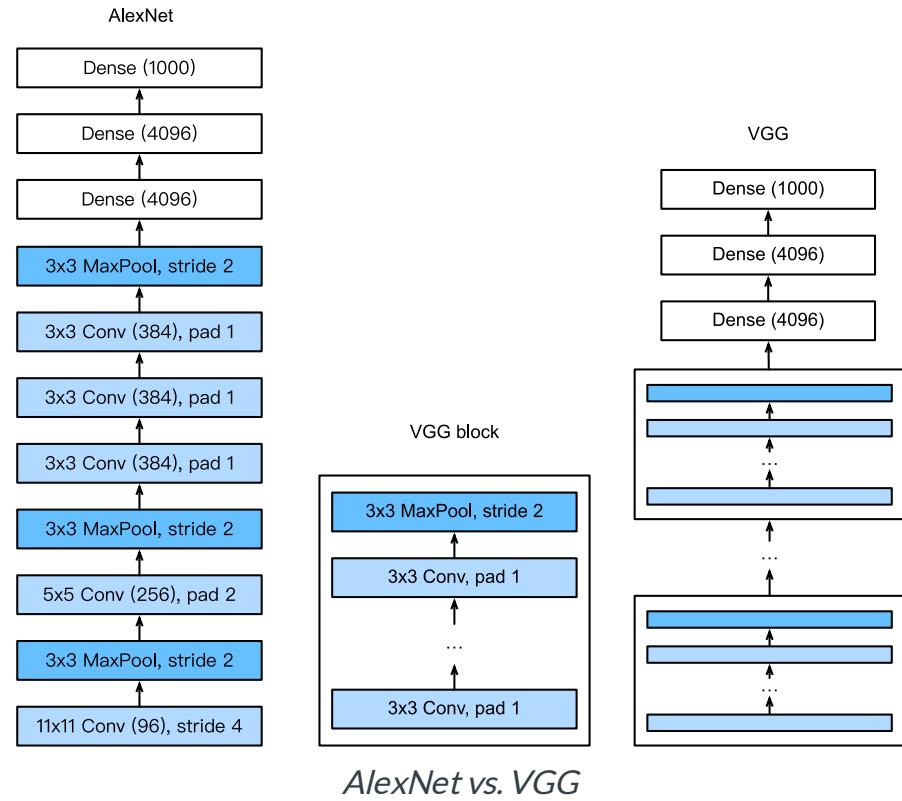


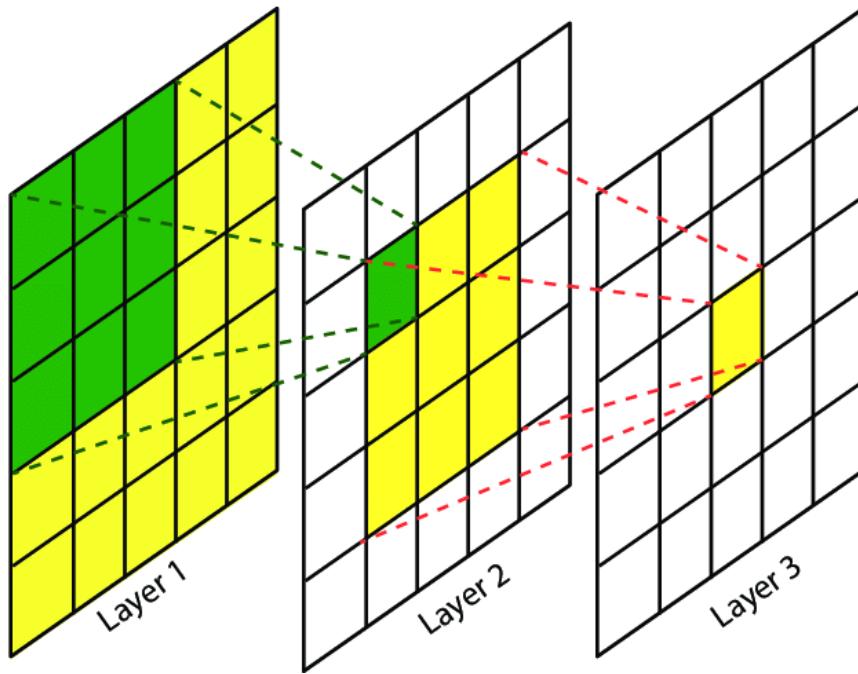
Layer (type)	Output Shape	Param #
<hr/>		
Conv2d-1	[-1, 64, 55, 55]	23,296
ReLU-2	[-1, 64, 55, 55]	0
MaxPool2d-3	[-1, 64, 27, 27]	0
Conv2d-4	[-1, 192, 27, 27]	307,392
ReLU-5	[-1, 192, 27, 27]	0
MaxPool2d-6	[-1, 192, 13, 13]	0
Conv2d-7	[-1, 384, 13, 13]	663,936
ReLU-8	[-1, 384, 13, 13]	0
Conv2d-9	[-1, 256, 13, 13]	884,992
ReLU-10	[-1, 256, 13, 13]	0
Conv2d-11	[-1, 256, 13, 13]	590,080
ReLU-12	[-1, 256, 13, 13]	0
MaxPool2d-13	[-1, 256, 6, 6]	0
Dropout-14	[-1, 9216]	0
Linear-15	[-1, 4096]	37,752,832
ReLU-16	[-1, 4096]	0
Dropout-17	[-1, 4096]	0
Linear-18	[-1, 4096]	16,781,312
ReLU-19	[-1, 4096]	0
Linear-20	[-1, 1000]	4,097,000
<hr/>		
Total params: 61,100,840		
Trainable params: 61,100,840		
Non-trainable params: 0		
<hr/>		
Input size (MB): 0.57		
Forward/backward pass size (MB): 8.31		
Params size (MB): 233.08		
Estimated Total Size (MB): 241.96		
<hr/>		

VGG (Simonyan and Zisserman, 2014)

Composition of 5 VGG blocks consisting of **CONV + POOL** layers, followed by a block of fully connected layers.

The network depth increased up to 19 layers, while the kernel sizes reduced to 3.





The **effective receptive field** is the part of the visual input that affects a given unit indirectly through previous convolutional layers. It grows linearly with depth.

E.g., a stack of two 3×3 kernels of stride 1 has the same effective receptive field as a single 5×5 kernel, but fewer parameters.

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 224, 224]	1,792
ReLU-2	[-1, 64, 224, 224]	0
Conv2d-3	[-1, 64, 224, 224]	36,928
ReLU-4	[-1, 64, 224, 224]	0
MaxPool2d-5	[-1, 64, 112, 112]	0
Conv2d-6	[-1, 128, 112, 112]	73,856
ReLU-7	[-1, 128, 112, 112]	0
Conv2d-8	[-1, 128, 112, 112]	147,584
ReLU-9	[-1, 128, 112, 112]	0
MaxPool2d-10	[-1, 128, 56, 56]	0
Conv2d-11	[-1, 256, 56, 56]	295,168
ReLU-12	[-1, 256, 56, 56]	0
Conv2d-13	[-1, 256, 56, 56]	590,080
ReLU-14	[-1, 256, 56, 56]	0
Conv2d-15	[-1, 256, 56, 56]	590,080
ReLU-16	[-1, 256, 56, 56]	0
MaxPool2d-17	[-1, 256, 28, 28]	0
Conv2d-18	[-1, 512, 28, 28]	1,180,160
ReLU-19	[-1, 512, 28, 28]	0
Conv2d-20	[-1, 512, 28, 28]	2,359,808
ReLU-21	[-1, 512, 28, 28]	0
Conv2d-22	[-1, 512, 28, 28]	2,359,808
ReLU-23	[-1, 512, 28, 28]	0
MaxPool2d-24	[-1, 512, 14, 14]	0
Conv2d-25	[-1, 512, 14, 14]	2,359,808
ReLU-26	[-1, 512, 14, 14]	0
Conv2d-27	[-1, 512, 14, 14]	2,359,808
ReLU-28	[-1, 512, 14, 14]	0
Conv2d-29	[-1, 512, 14, 14]	2,359,808
ReLU-30	[-1, 512, 14, 14]	0
MaxPool2d-31	[-1, 512, 7, 7]	0
Linear-32	[-1, 4096]	102,764,544
ReLU-33	[-1, 4096]	0
Dropout-34	[-1, 4096]	0
Linear-35	[-1, 4096]	16,781,312
ReLU-36	[-1, 4096]	0
Dropout-37	[-1, 4096]	0
Linear-38	[-1, 1000]	4,097,000

Total params: 138,357,544

Trainable params: 138,357,544

Non-trainable params: 0

Input size (MB): 0.57

Forward/backward pass size (MB): 218.59

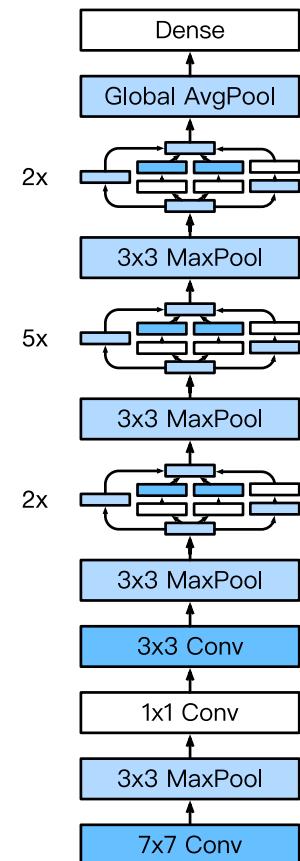
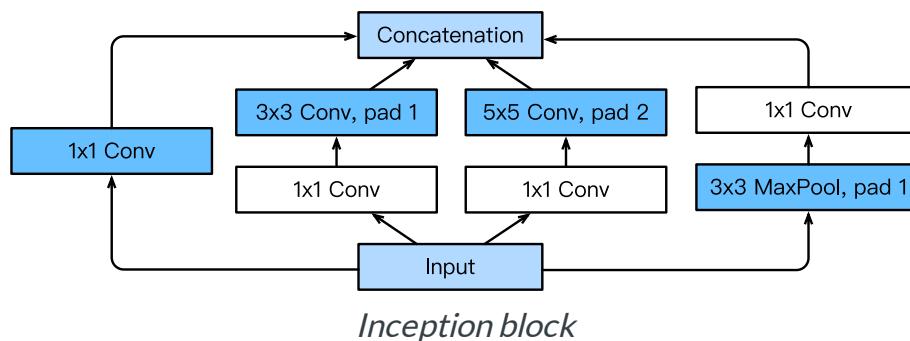
Params size (MB): 527.79

Estimated Total Size (MB): 746.96

GoogLeNet (Szegedy et al, 2014)

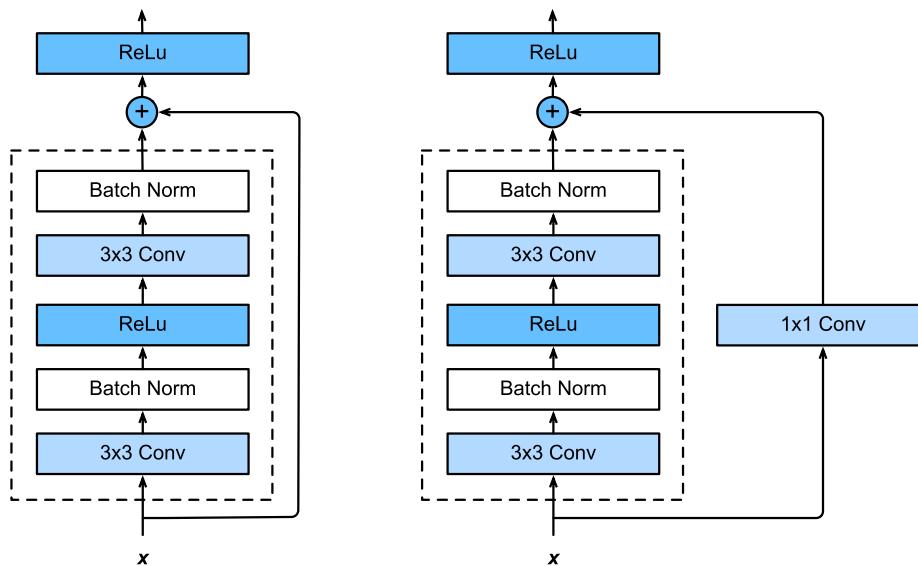
Composition of two **CONV + POOL** layers, a stack of 9 inception blocks, and a global average pooling layer.

Each inception block is itself defined as a convolutional network with 4 parallel paths.

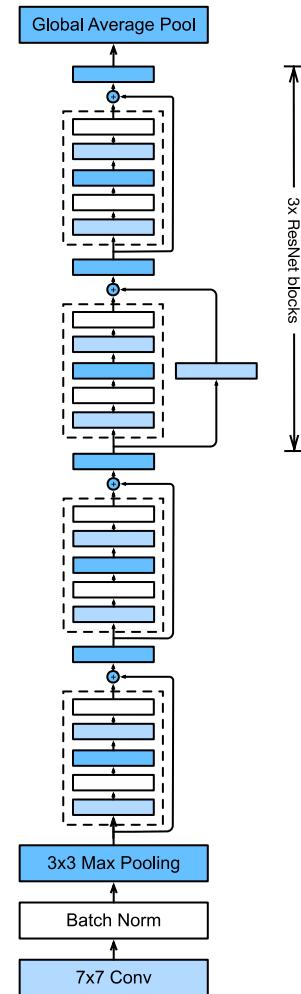


ResNet (He et al, 2015)

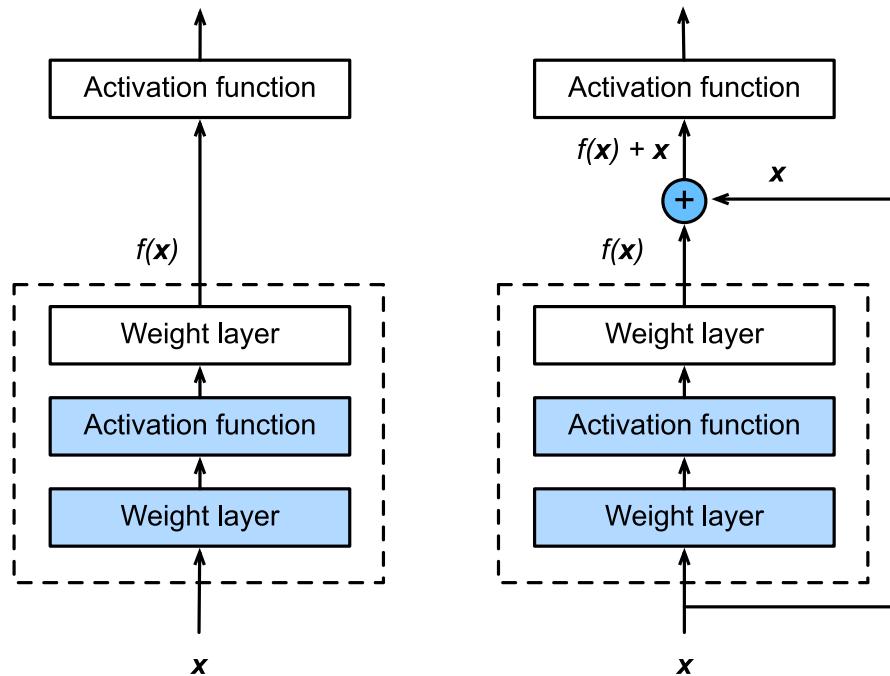
Composition of first layers similar to GoogLeNet, a stack of 4 residual blocks, and a global average pooling layer. Extensions consider more residual blocks, up to a total of 152 layers (ResNet-152).



Regular ResNet block vs. ResNet block with 1×1 convolution.

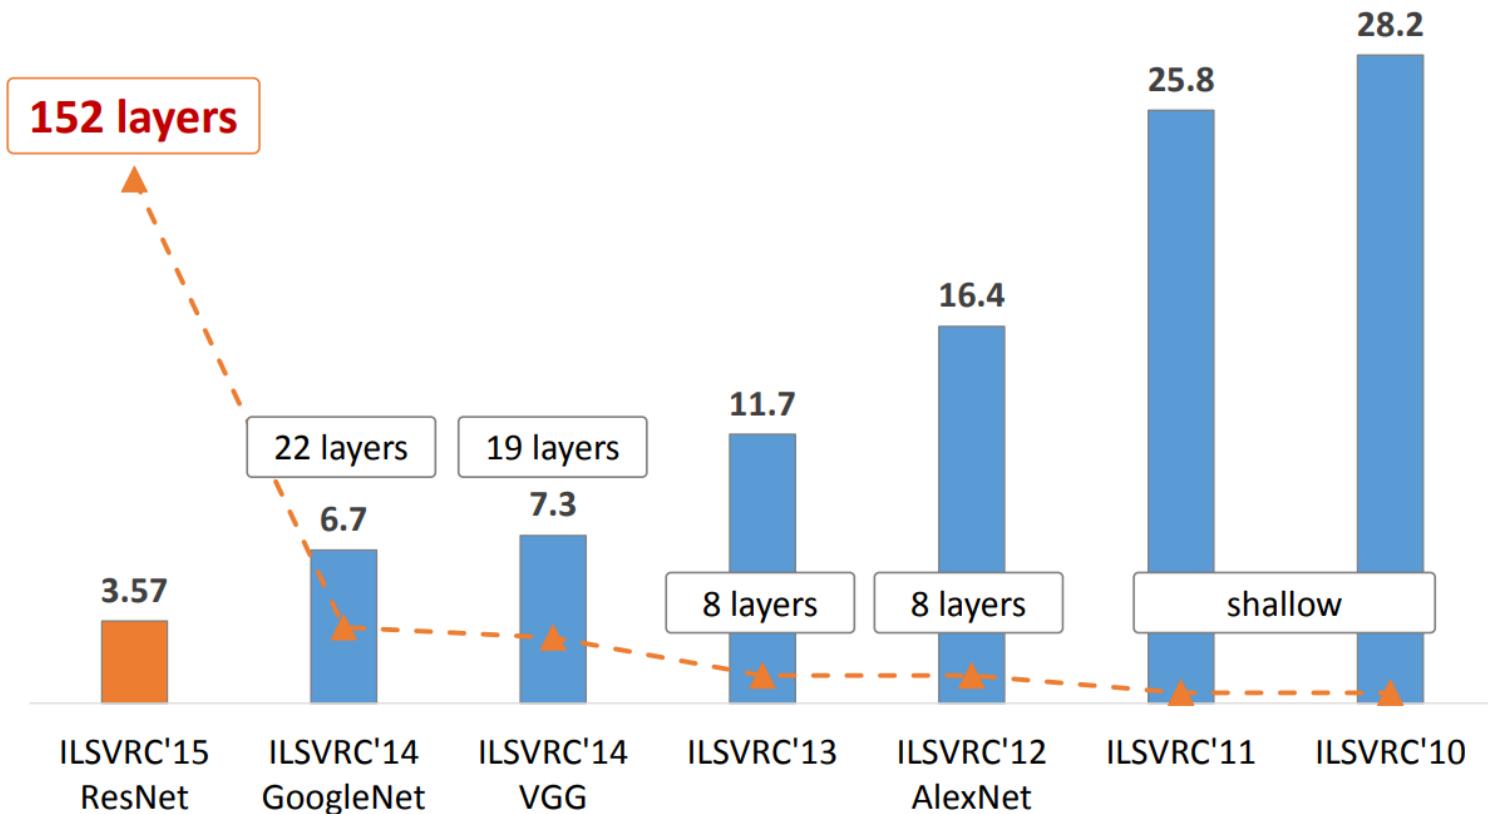


Training networks of this depth is made possible because of the **skip connections** in the residual blocks. They allow the gradients to shortcut the layers and pass through without vanishing.



Layer (type)	Output Shape	Param #	...
Conv2d-1	[−1, 64, 112, 112]	9,408	Bottleneck-130
BatchNorm2d-2	[−1, 64, 112, 112]	128	Conv2d-131
ReLU-3	[−1, 64, 112, 112]	0	BatchNorm2d-132
MaxPool2d-4	[−1, 64, 56, 56]	0	ReLU-133
Conv2d-5	[−1, 64, 56, 56]	4,096	Conv2d-134
BatchNorm2d-6	[−1, 64, 56, 56]	128	BatchNorm2d-135
ReLU-7	[−1, 64, 56, 56]	0	ReLU-136
Conv2d-8	[−1, 64, 56, 56]	36,864	Conv2d-137
BatchNorm2d-9	[−1, 64, 56, 56]	128	BatchNorm2d-138
ReLU-10	[−1, 64, 56, 56]	0	ReLU-139
Conv2d-11	[−1, 256, 56, 56]	16,384	Bottleneck-140
BatchNorm2d-12	[−1, 256, 56, 56]	512	Conv2d-141
Conv2d-13	[−1, 256, 56, 56]	16,384	BatchNorm2d-142
BatchNorm2d-14	[−1, 256, 56, 56]	512	ReLU-143
ReLU-15	[−1, 256, 56, 56]	0	Conv2d-144
Bottleneck-16	[−1, 256, 56, 56]	0	BatchNorm2d-145
Conv2d-17	[−1, 64, 56, 56]	16,384	ReLU-146
BatchNorm2d-18	[−1, 64, 56, 56]	128	Conv2d-147
ReLU-19	[−1, 64, 56, 56]	0	BatchNorm2d-148
Conv2d-20	[−1, 64, 56, 56]	36,864	Conv2d-149
BatchNorm2d-21	[−1, 64, 56, 56]	128	BatchNorm2d-150
ReLU-22	[−1, 64, 56, 56]	0	ReLU-151
Conv2d-23	[−1, 256, 56, 56]	16,384	Bottleneck-152
BatchNorm2d-24	[−1, 256, 56, 56]	512	Conv2d-153
ReLU-25	[−1, 256, 56, 56]	0	BatchNorm2d-154
Bottleneck-26	[−1, 256, 56, 56]	0	ReLU-155
Conv2d-27	[−1, 64, 56, 56]	16,384	Conv2d-156
BatchNorm2d-28	[−1, 64, 56, 56]	128	BatchNorm2d-157
ReLU-29	[−1, 64, 56, 56]	0	ReLU-158
Conv2d-30	[−1, 64, 56, 56]	36,864	Conv2d-159
BatchNorm2d-31	[−1, 64, 56, 56]	128	BatchNorm2d-160
ReLU-32	[−1, 64, 56, 56]	0	ReLU-161
Conv2d-33	[−1, 256, 56, 56]	16,384	Bottleneck-162
BatchNorm2d-34	[−1, 256, 56, 56]	512	Conv2d-163
ReLU-35	[−1, 256, 56, 56]	0	BatchNorm2d-164
Bottleneck-36	[−1, 256, 56, 56]	0	ReLU-165
Conv2d-37	[−1, 128, 56, 56]	32,768	Conv2d-166
BatchNorm2d-38	[−1, 128, 56, 56]	256	BatchNorm2d-167
ReLU-39	[−1, 128, 56, 56]	0	ReLU-168
Conv2d-40	[−1, 128, 28, 28]	147,456	Conv2d-169
BatchNorm2d-41	[−1, 128, 28, 28]	256	BatchNorm2d-170
ReLU-42	[−1, 128, 28, 28]	0	ReLU-171
Conv2d-43	[−1, 512, 28, 28]	65,536	Bottleneck-172
BatchNorm2d-44	[−1, 512, 28, 28]	1,024	AvgPool2d-173
Conv2d-45	[−1, 512, 28, 28]	131,072	Linear-174
BatchNorm2d-46	[−1, 512, 28, 28]	1,024	Total params: 25,557,032
ReLU-47	[−1, 512, 28, 28]	0	Trainable params: 25,557,032
Bottleneck-48	[−1, 512, 28, 28]	0	Non-trainable params: 0
Conv2d-49	[−1, 128, 28, 28]	65,536	Input size (MB): 0.57
BatchNorm2d-50	[−1, 128, 28, 28]	256	Forward/backward pass size (MB): 286.56
ReLU-51	[−1, 128, 28, 28]	0	Params size (MB): 97.49
Conv2d-52	[−1, 128, 28, 28]	147,456	Estimated Total Size (MB): 384.62
BatchNorm2d-53	[−1, 128, 28, 28]	256	-----

The benefits of depth



Under the hood

Understanding what is happening in deep neural networks after training is complex and the tools we have are limited.

In the case of convolutional neural networks, we can look at:

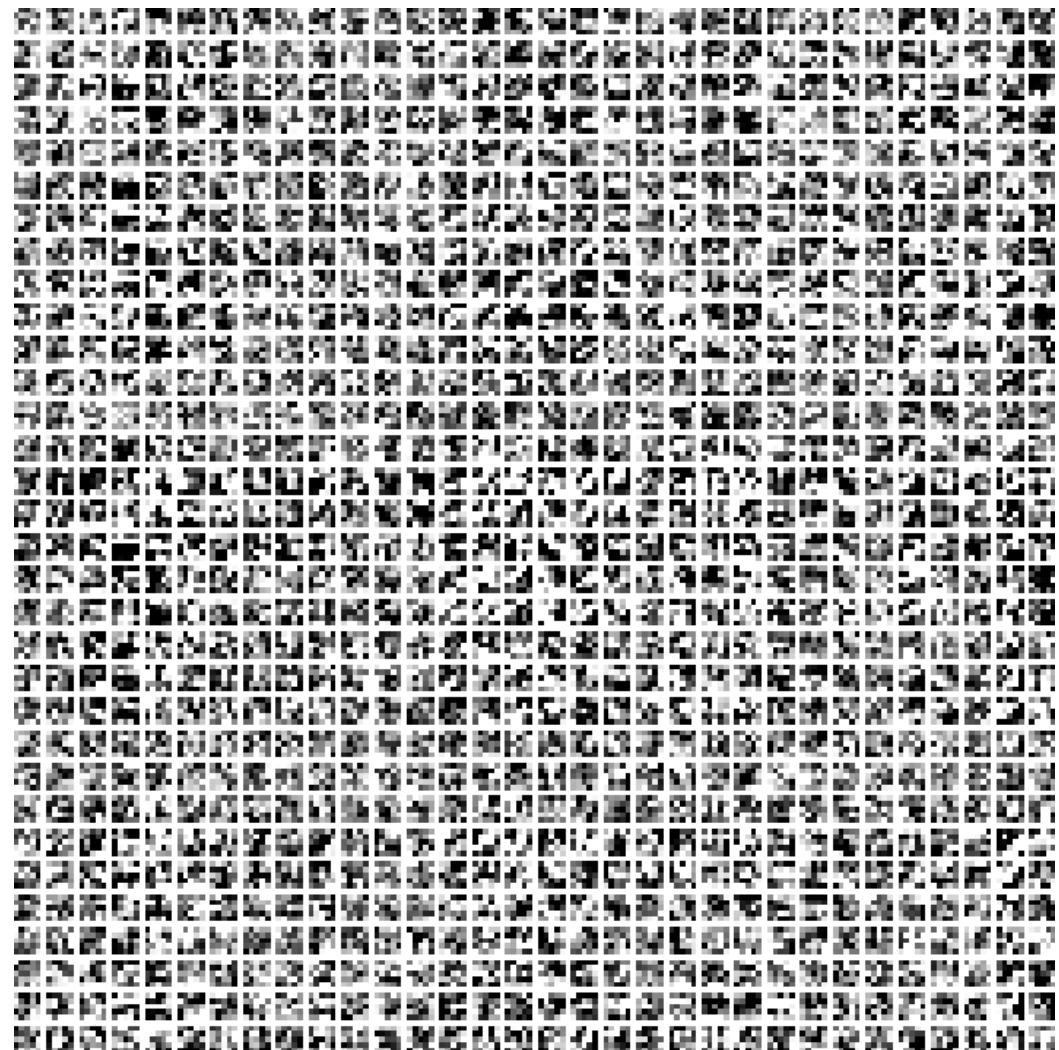
- the network's kernels as images
- internal activations on a single sample as images
- distributions of activations on a population of samples
- derivatives of the response with respect to the input
- maximum-response synthetic samples

Looking at filters

LeNet's first convolutional layer, all filters.

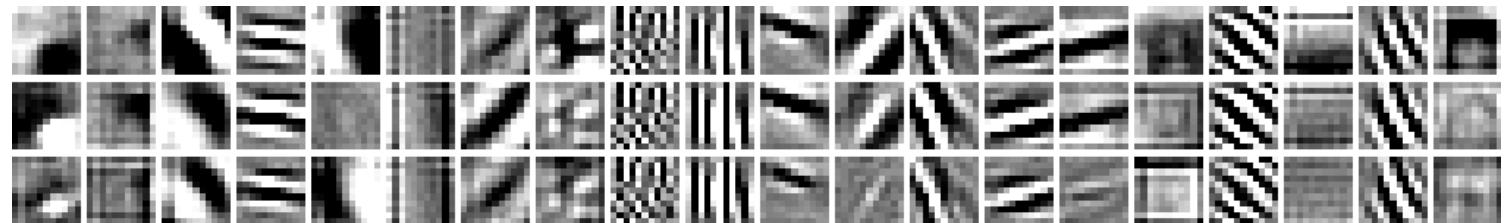


LeNet's second convolutional layer, first 32 filters.



The image shows a 28x28 grid of 32 handwritten digit filters. These filters are the output of the second convolutional layer of the LeNet architecture. Each filter is a 5x5 kernel applied to the 28x28 input image. The filters capture various features such as edges, corners, and specific patterns characteristic of digits like '4', '7', and '2'. The filters are arranged in a grid where each row contains 8 filters and each column contains 4 filters.

AlexNet's first convolutional layer, first 20 filters.

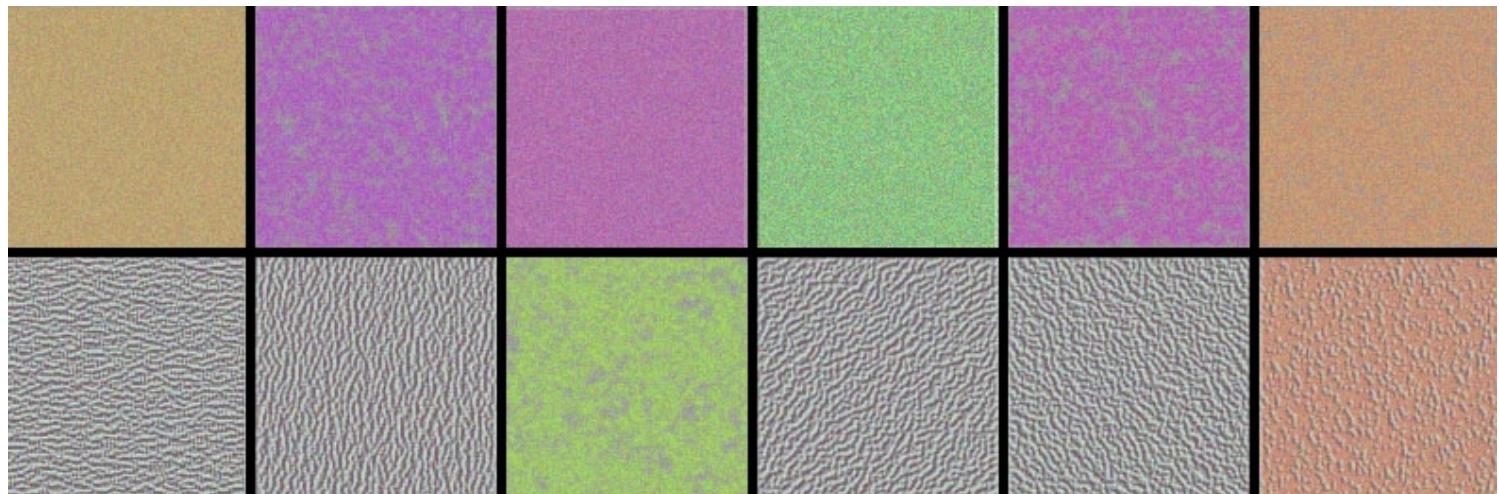


Maximum response samples

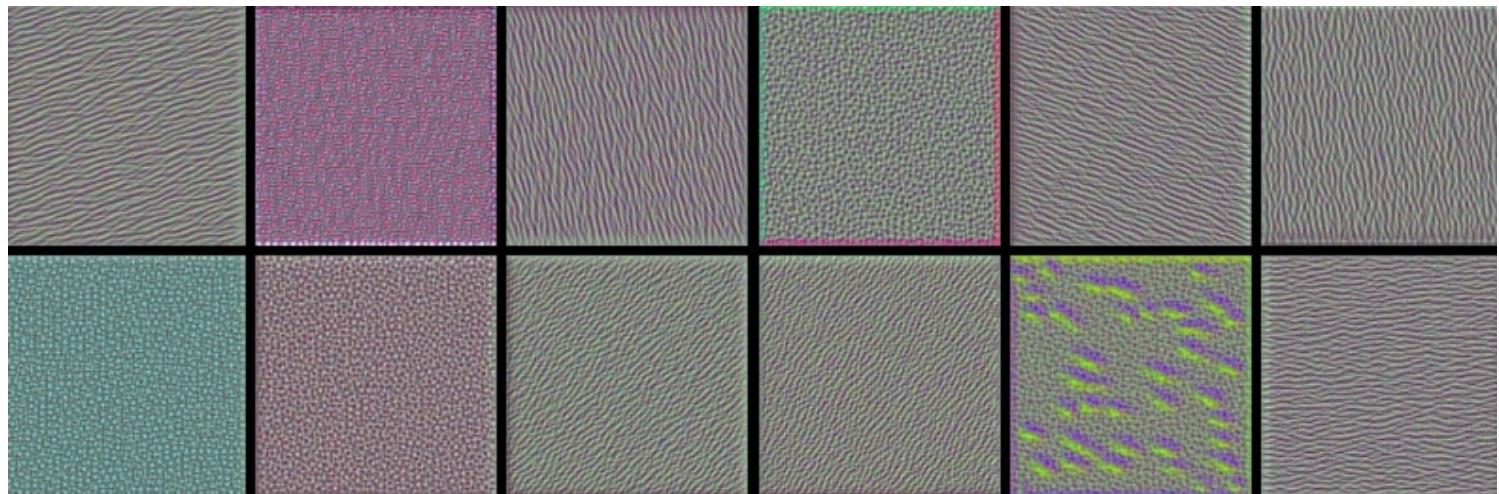
Convolutional networks can be inspected by looking for synthetic input images \mathbf{x} that maximize the activation $\mathbf{h}_{\ell,d}(\mathbf{x})$ of a chosen convolutional kernel \mathbf{u} at layer ℓ and index d in the layer filter bank.

These samples can be found by gradient ascent on the input space:

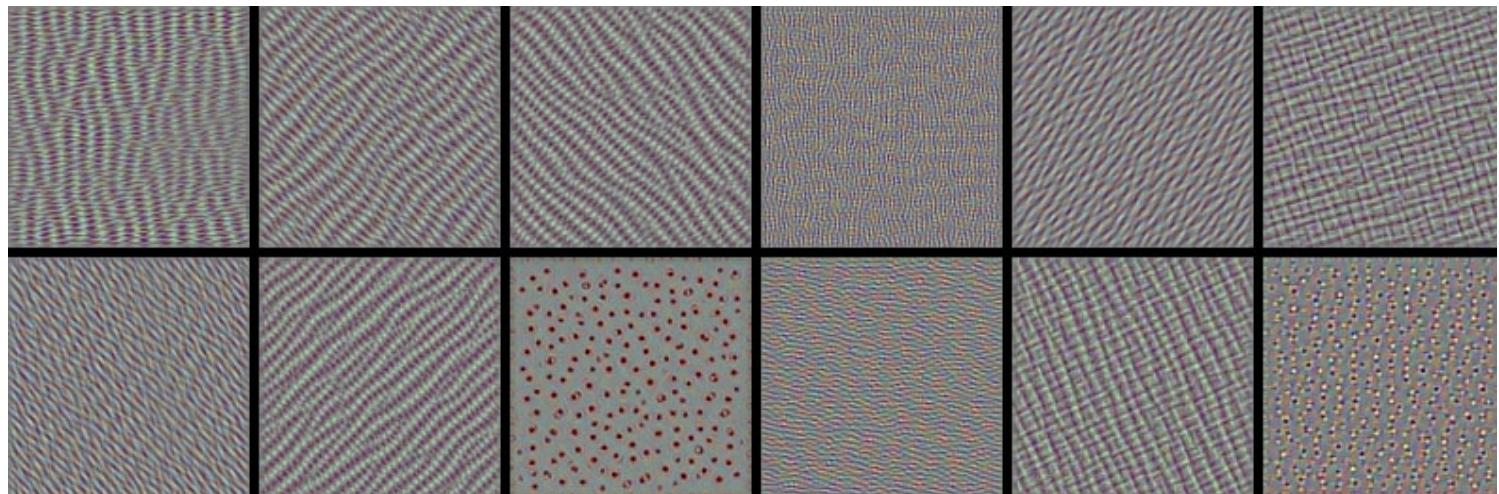
$$\begin{aligned}\mathcal{L}_{\ell,d}(\mathbf{x}) &= \|\mathbf{h}_{\ell,d}(\mathbf{x})\|_2 \\ \mathbf{x}_0 &\sim U[0, 1]^{C \times H \times W} \\ \mathbf{x}_{t+1} &= \mathbf{x}_t + \gamma \nabla_{\mathbf{x}} \mathcal{L}_{\ell,d}(\mathbf{x}_t)\end{aligned}$$



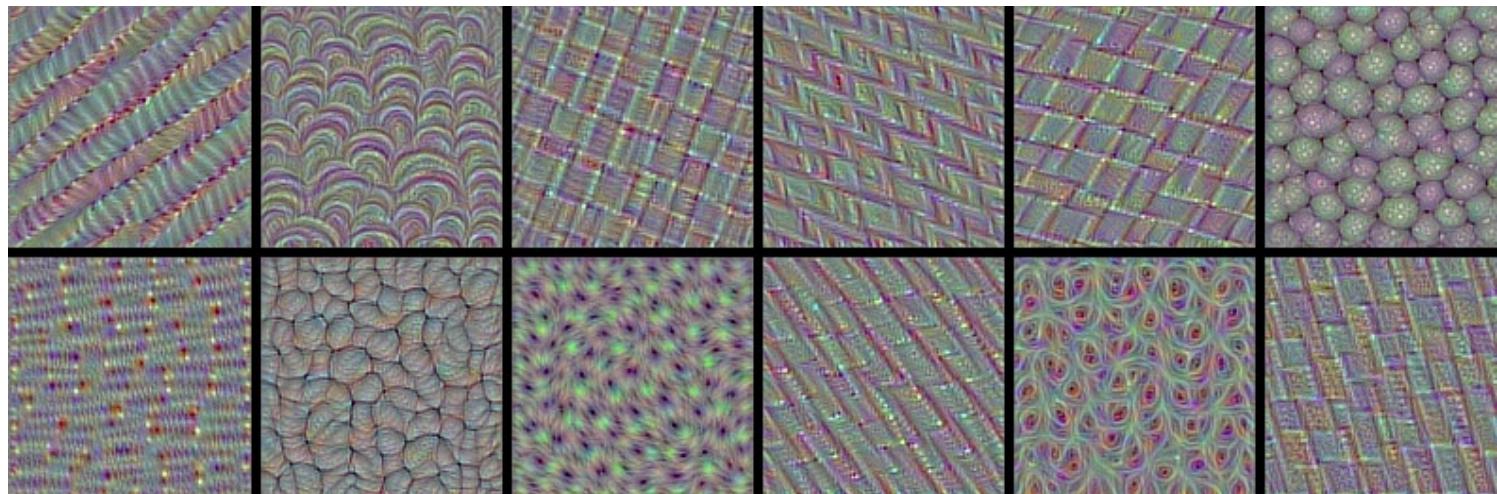
VGG-16, convolutional layer 1-1, a few of the 64 filters



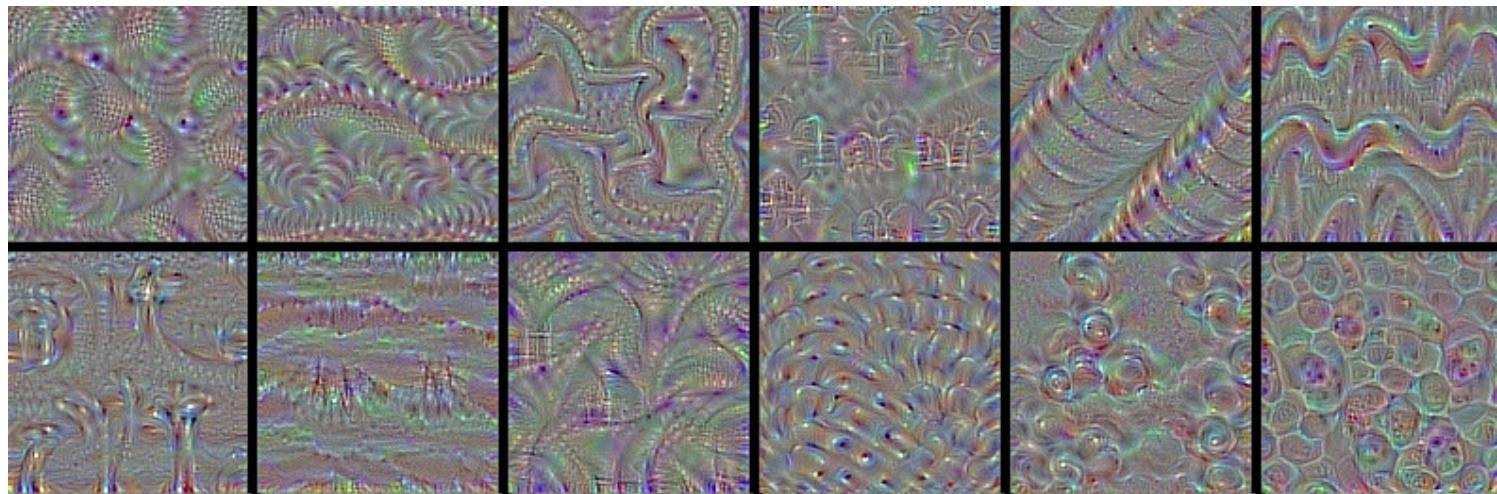
VGG-16, convolutional layer 2-1, a few of the 128 filters



VGG-16, convolutional layer 3-1, a few of the 256 filters



VGG-16, convolutional layer 4-1, a few of the 512 filters

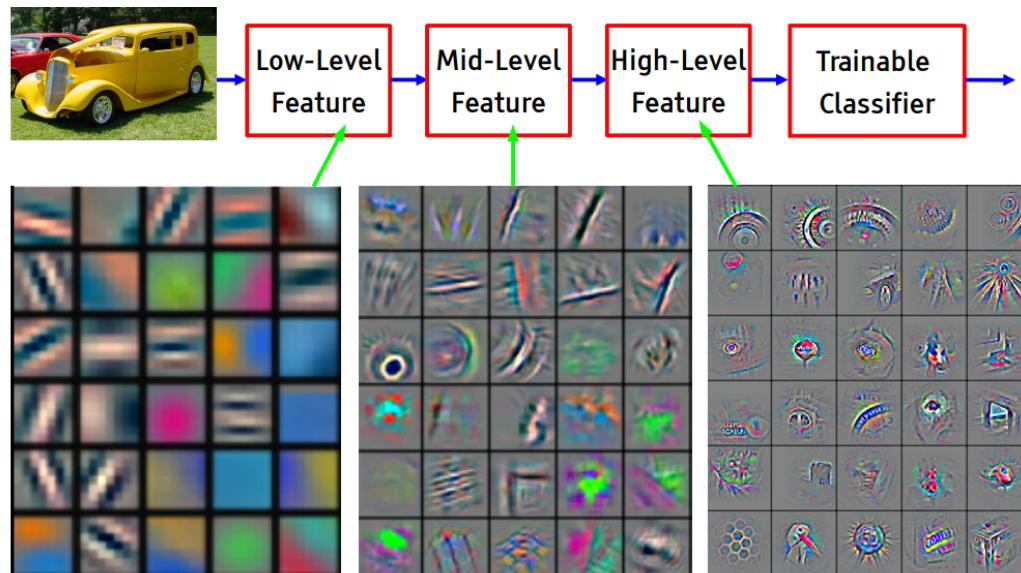


VGG-16, convolutional layer 5-1, a few of the 512 filters

Some observations:

- The first layers appear to encode direction and color.
- The direction and color filters get combined into grid and spot textures.
- These textures gradually get combined into increasingly complex patterns.

The network appears to learn a **hierarchical composition of patterns**.

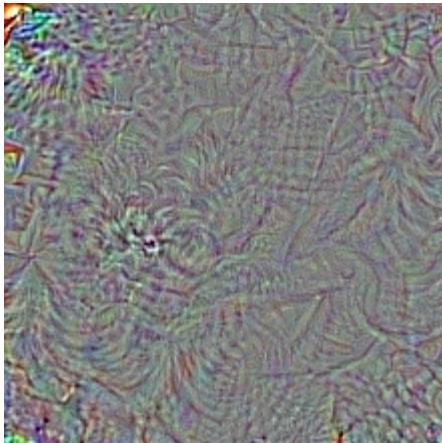


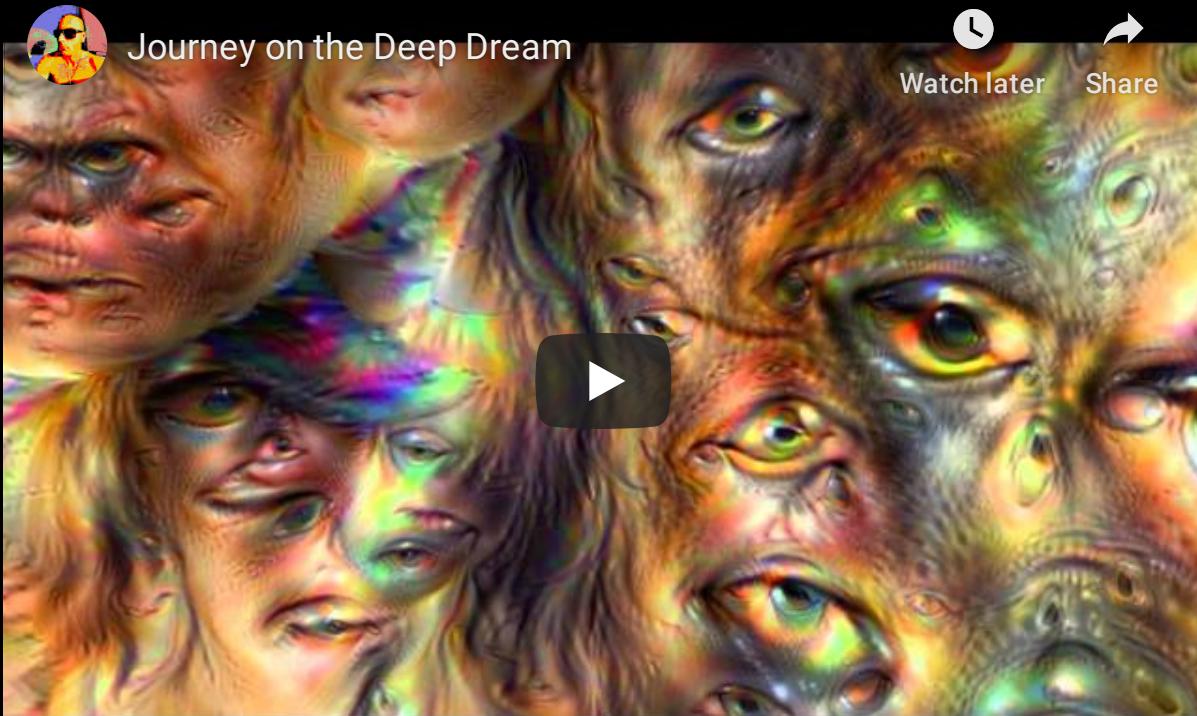
Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

What if we build images that maximize the activation of a chosen class output?

What if we build images that maximize the activation of a chosen class output?

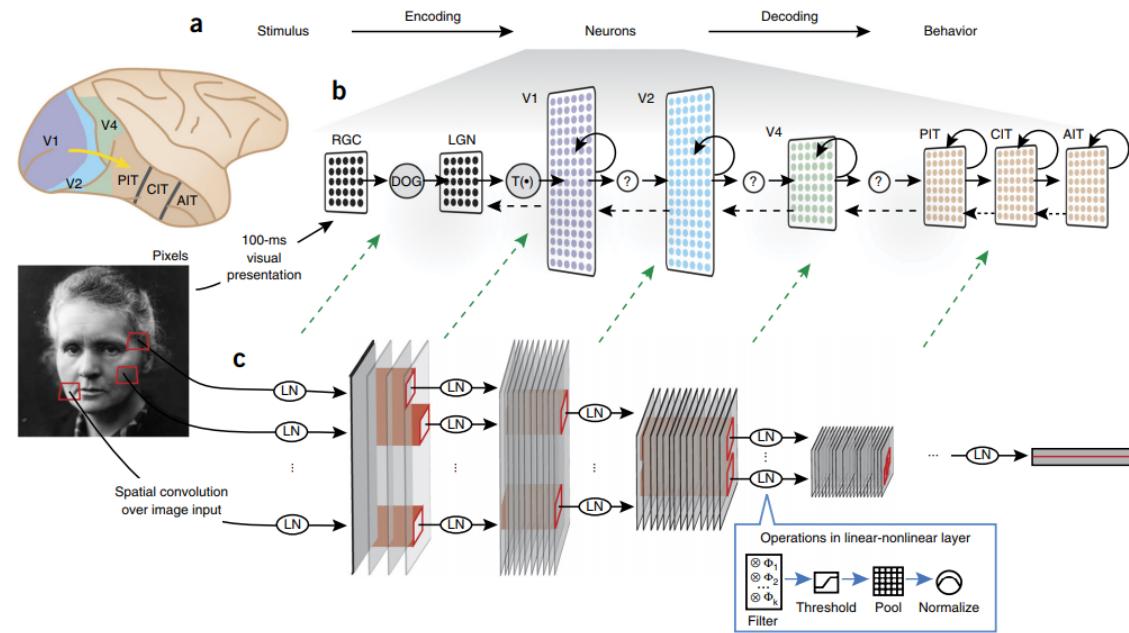
The left image is predicted **with 99.9% confidence** as a magpie!





Deep Dream. Start from an image \mathbf{x}_t , offset by a random jitter, enhance some layer activation at multiple scales, zoom in, repeat on the produced image \mathbf{x}_{t+1} .

Biological plausibility



"Deep hierarchical neural networks are beginning to transform neuroscientists' ability to produce quantitatively accurate computational models of the sensory systems, especially in higher cortical areas where neural response properties had previously been enigmatic."

The end.

References

- Francois Fleuret, Deep Learning Course, [4.4. Convolutions](#), EPFL, 2018.
- Yannis Avrithis, Deep Learning for Vision, [Lecture 1: Introduction](#), University of Rennes 1, 2018.
- Yannis Avrithis, Deep Learning for Vision, [Lecture 7: Convolution and network architectures](#), University of Rennes 1, 2018.
- Olivier Grisel and Charles Ollion, Deep Learning, [Lecture 4: Convolutional Neural Networks for Image Classification](#), Université Paris-Saclay, 2018.